

~~X~~~~4/27~~

## II Structures

— INSERT B goes here —

Now that we understand how to construct a good proof, we can turn our attention to

the structures that are most useful in modelling problems that arise in computer science.

\* *in computer science*

The most important structure is a *graph* (aka *network*). Graphs provide an excellent

mechanism for modelling associations between pairs of objects (e.g., two exams that

## INSERT B

Structure is fundamental in computer science. Whether you are writing code, solving an optimization problem, or designing a network, you will be dealing with structure. The better you ~~can~~ can understand the structure, and ~~the better~~ the better your results will be. And if you can reason about structure, then you will be in a good position to convince others (and yourself) that your results are worthy.

cannot be given at the same time, two people that like each other, or two subroutines

that can be run independently). In Chapter 5, we study graphs that represent *symmetric*

*like those just mentioned. In*

relationships, ~~and~~ in Chapter 6, we consider graphs where the relationship is *one-way*

(e.g., a situation where you can go from  $x$  to  $y$  but not necessarily vice-versa.)

In Chapter 7, we consider the more general notion of a *relation* and we examine important classes of relations such as partially ordered sets. Partially ordered sets arise frequently in scheduling problems.

We conclude in Chapter 8 with a study of *state machines*. State machines can be used to model a variety of processes and are a fundamental tool in proving that an algorithm terminates and that it produces the correct output.

X

Fonts are too small  
- it seems they got smaller, but  
maybe it's just my eyes getting  
weaker.

## 6

## Directed graphs

---

### 6.1 Definitions

So far, we have been working with graphs with undirected edges. A *directed edge* is an edge

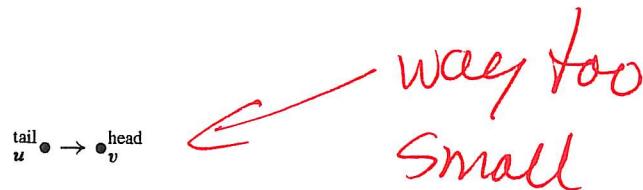
where the endpoints are distinguished—one is the *head* and one is the *tail*. In particular, a

directed edge is specified as an ordered pair of vertices  $u, v$  and is denoted by  $(u, v)$  or  $u \rightarrow v$ .

X

In this case,  $u$  is the tail of the edge and  $v$  is the head. For example, see Figure 6.1.

A graph with directed edges is called a *directed graph* or *digraph*.



Redraft graphic: Missing graphic

Figure 6.1: A directed edge  $e = (u, v)$ .  $u$  is the tail of  $e$  and  $v$  is the head of  $e$ .

**Definition 6.1.1.** A directed graph  $G = (V, E)$  consists of a nonempty set of nodes  $V$  and a set

of directed edges  $E$ . Each edge  $e$  of  $E$  is specified by an ordered pair of vertices  $u, v \in V$  where ~~where  $u \neq v$~~

~~$u \neq v$~~  A directed graph is *simple* if it has no *loops* (i.e., edges of the form  $u \rightarrow u$ ) [Didn't you just say that couldn't be an edge?] and no multiple edges.

Since we will focus on the case of simple directed graphs in this chapter, we will generally

omit the word *simple* when referring to them. Note that such a graph can contain ~~one~~ <sup>an</sup> edge  $u \rightarrow v$

as well as the edge  $v \rightarrow u$  since these are different edges (that is, they have a different tail).

Directed graphs arise in applications where the relationship represented by an edge is 1-way

or asymmetric. Examples include: a 1-way street, one person likes another but the feeling *isn't*

not necessarily reciprocated, a communication channel such as a cable modem that has more

capacity for downloading than uploading, one entity is larger than another, and one job needs to

be completed before another can begin. We'll see several such examples in this chapter and also

in Chapter 7.

Most all of the definitions for undirected graphs from Chapter 5 carry over in a natural way

for directed graphs. For example, two directed graphs  $G_1 = (V_1, E_1)$  and  $G_2 = (V_2, E_2)$  are

*isomorphic* if there exists a bijection  $f : V_1 \rightarrow V_2$  such that for every pair of vertices  $u, v \in V_1$ ,

$$u \rightarrow v \in E_1 \quad \text{IFF} \quad f(u) \rightarrow f(v) \in E_2.$$

Directed graphs have adjacency matrices just like undirected graphs. In the case of a directed

graph  $G = (V, E)$ , the adjacency matrix  $A_G = \{a_{ij}\}$  is defined so that

$$a_{ij} = \begin{cases} 1 & \text{if } i \rightarrow j \in E \\ 0 & \text{otherwise.} \end{cases}$$

Redraft graphic: Missing graphic

Figure 6.2: A 4-node directed graph with 6 edges.

The only difference is that the adjacency matrix for a directed graph is not necessarily symmetric (that is, it may be that  $A_G^T \neq A_G$ ).

### 6.1.1 Degrees

With directed graphs, the notion of degree splits into *indegree* and *outdegree*. For example,

$\text{indegree}(c) = 2$  and  $\text{outdegree}(c) = 1$  for the graph in Figure 6.2. If a node has outdegree 0, it is called a *sink*; if it has indegree 0, it is called a *source*. The graph in Figure 6.2 has one source (node  $a$ ) and no sinks.

### 6.1.2 Directed Walks, Paths, and Cycles

The definitions for (directed) walks, paths, and cycles in a directed graph are similar to those for undirected graphs except that the direction of the edges need to be consistent with the order in which the walk is traversed.

**Definition 6.1.2.** A *directed walk* (or more simply, a *walk*) in a directed graph  $G$  is a sequence of vertices  $v_0, v_1, \dots, v_k$  and edges

$$v_0 \rightarrow v_1, v_1 \rightarrow v_2, \dots, v_{k-1} \rightarrow v_k$$

such that  $v_{i-1} \rightarrow v_i$  is an edge of  $G$  for all  $i$  where  $0 \leq i < k$ . A *directed path* (or *path*) in a directed graph is a walk where the nodes in the walk are all different. A *directed closed walk* (or *closed walk*) in a directed graph is a walk where  $v_0 = v_k$ . A *directed cycle* (or *cycle*) in a directed graph is a closed walk where all the vertices  $v_i$  are different for  $0 \leq i < k$ .

As with undirected graphs, we will typically refer to a walk in a directed graph by a sequence

of vertices. For example, for the graph in Figure 6.2,

- $a, b, c, b, d$  is a walk,
- $a, b, d$  is a path,
- $d, c, b, c, b, d$  is a closed walk, and
- $b, d, c, b$  is a cycle.

Note that  $b, c, b$  is also a cycle for the graph in Figure 6.2. This is a cycle of length 2. Such cycles are not possible with undirected graphs.

Also note that

$$c, b, a, d$$

is *not* a walk in the graph shown in Figure 6.2, since  $b \rightarrow a$  is not an edge in this graph. (You are *not* allowed to traverse edges in the wrong direction as part of a walk.)

X

A path or cycle in a directed graph is said to be *Hamiltonian* if it visits every node in the graph.

For example,  $a, b, d, c$  is the only Hamiltonian path for the graph in Figure 6.2. The graph in Figure 6.2 does not have a Hamiltonian cycle.

A walk in a directed graph is said to be *Eulerian* if it contains every edge. The graph shown in Figure 6.2 does not have an Eulerian walk. Can you see why not? (Hint: Look at node  $a$ .)

### 6.1.3 Strong Connectivity

The notion of being connected is a little more complicated for a directed graph than it is for an undirected graph. For example, should we consider the graph in Figure 6.2 to be connected?

There is a path from node  $a$  to every other node so on that basis, we might answer “Yes.” But there is ~~not~~ path from nodes  $b, c$ , or  $d$  to node  $a$ , and so on that basis, we might answer “No.”

For this reason, graph theorists have come up with the notion of strong connectivity for directed graphs.



X

X



**Definition 6.1.3.** A directed graph  $G = (V, E)$  is said to be *strongly connected* if for every pair of nodes  $u, v \in V$ , there is a directed path from  $u$  to  $v$  (and vice-versa) in  $G$ .

For example, the graph in Figure 6.2 is not strongly connected since there is no directed path from node  $b$  to node  $a$ . But if node  $a$  is removed, the resulting graph would be strongly connected.

A directed graph is said to be *weakly connected* (or, more simply, *connected*) if the corresponding undirected graph (where directed edges  $u \rightarrow v$  and/or  $v \rightarrow u$  are replaced with a single undirected edge  $u \xrightarrow{\text{---}} v$ ) is connected. For example, the graph in Figure 6.2 is weakly connected.



#### 6.1.4 DAGs

If an undirected graph does not have any cycles, then it is a tree or a forest. But what does a directed graph look like if it has no cycles? For example, consider the graph in Figure 6.3. This

X

Redraft graphic: Missing graphic

Figure 6.3: A 4-node directed acyclic graph (DAG).

graph is weakly connected and has no directed cycles but it certainly does not look like a tree.

**Definition 6.1.4.** A directed graph is called a *directed acyclic graph* (or, *DAG*) if it does not contain any directed cycles.

A first glance, DAGs don't appear to be particularly interesting. But first impressions are not always accurate. In fact, DAGs arise in many scheduling and optimization problems and they have several interesting properties. We will study them extensively in Chapter 7.

---

## 6.2 Tournament Graphs

and that

Suppose that  $n$  players compete in a round-robin tournament. Thus for every pair of players  $u$  and  $v$ , either  $u$  beats  $v$  or  $v$  beats  $u$ . Interpreting the results of a round-robin tournament can be

X

Redraft graphic: Missing graphic

Figure 6.4: A 5-node tournament graph.

problematic. There might be all sorts of cycles where  $x$  beat  $y$ ,  $y$  beat  $z$ , yet  $z$  beat  $x$ . Graph

~~does not solve this problem but it can provide some interesting perspectives.~~

and who is the best player?

~~theory provides at least a partial solution to this problem.~~

The results of a round-robin tournament can be represented with a *tournament graph*. This is

a directed graph in which the vertices represent players and the edges indicate the outcomes of

games. In particular, an edge from  $u$  to  $v$  indicates that player  $u$  defeated player  $v$ . In a round-

robin tournament, every pair of players has a match. Thus, in a tournament graph there is either

an edge from  $u$  to  $v$  or an edge from  $v$  to  $u$  (but not both) for *every* pair of distinct vertices  $u$

and  $v$ . An example of a tournament graph is shown in Figure 6.4.

RFB

~~Figure~~

### 6.2.1 Finding a Hamiltonian Path in a Tournament Graph

We're going to prove that in every round-robin tournament, there exists a ranking of the players

such that each player lost to the player one position higher. For example, in the tournament corresponding to Figure 6.4, the ranking

$$a > b > d > e > c$$

satisfies this criterion, because  $b$  lost to  $a$ ,  $d$  lost to  $b$ ,  $e$  lost to  $d$ , and  $c$  lost to  $e$ . In graph terms, proving the existence of such a ranking amounts to proving that every tournament graph has a Hamiltonian path.

**Theorem 6.2.1.** *Every tournament graph contains a directed Hamiltonian path.*

*Proof.* We use strong induction. Let  $P(n)$  be the proposition that every tournament graph with  $n$  vertices contains a directed Hamiltonian path.

**Base case:**  $P(1)$  is trivially true; every graph with a single vertex has a Hamiltonian path con-

**Redraft graphic: Missing graphic**

Figure 6.5: The sets  $T$  and  $F$  in a tournament graph.

sisting of only that vertex.

**Inductive step:** For  $n \geq 1$ , we assume that  $P(1), \dots, P(n)$  are all true and prove  $P(n + 1)$ .

Consider a tournament graph  $G = (V, E)$  with  $n + 1$  players. Select one vertex  $v$  arbitrarily. Ev-

ery other vertex in the tournament either has an edge *to* vertex  $v$  or an edge *from* vertex  $v$ . Thus,

we can partition the remaining vertices into two corresponding sets,  $T$  and  $F$ , each containing at

most  $n$  vertices, where  $T = \{u \mid u \rightarrow v \in E\}$  and  $F = \{u \mid v \rightarrow u \in E\}$ . For example, see

Figure 6.5.

The vertices in  $T$  together with the edges that join them form a smaller tournament. Thus, by

strong induction, there is a Hamiltonian path within  $T$ . Similarly, there is a Hamiltonian path

within the tournament on the vertices in  $F$ . Joining the path in  $T$  to the vertex  $v$  followed by the

path in  $F$  gives a Hamiltonian path through the whole tournament. As special cases, if  $T$  or  $F$  is empty, then so is the corresponding portion of the path. ■

The ranking defined by a Hamiltonian path is not entirely satisfactory. For example, in the tournament associated with Figure 6.4, notice that the lowest-ranked player,  $c$ , actually defeated the highest-ranked player,  $a$ .

In practice, players ~~are~~<sup>are</sup> typically ranked according to how many victories they achieve. This

makes sense for several reasons. One not-so-obvious reason is that if the player with the most victories does not beat some other player  $v$ , he is guaranteed to have at least beaten a third player

*We'll prove this fact shortly, but first, let's who beat  $v$ . We prove this fact in the next section.*

*talk about chickens.*

### 6.2.2 The King Chicken Theorem

*chickens are rather aggressive birds and tend to establish dominance relationships by pecking.*

Suppose that there are  $n$  chickens in a farmyard. For each pair of distinct chickens, either the

first pecks the second or the second pecks the first, but not both. We say that chicken  $u$  *virtually*

X

**Redraft graphic: Missing graphic**

Figure 6.6: A 4-chicken tournament in which chickens  $a$ ,  $b$ , and  $d$  are kings.

pecks chicken  $v$  if either:

~~either~~ directly

- Chicken  $u$  pecks chicken  $v$ , or
- Chicken  $u$  pecks some other chicken  $w$  who in turn pecks chicken  $v$ .

A chicken that virtually pecks every other chicken is called a *king chicken*.

We can model this situation with a tournament digraph. The vertices are chickens, and an edge

$u \rightarrow v$  indicates that chicken  $u$  pecks chicken  $v$ . In the tournament shown in Figure 6.6, three

of the four chickens are kings. Chicken  $c$  is not a king in this example since it does not peck

chicken  $b$  and it does not peck any chicken that pecks chicken  $b$ . Chicken  $a$  is a king since it

pecks chicken  $d$ , who in turn pecks chickens  $b$  and  $c$ .

**Theorem 6.2.2 (King Chicken Theorem).** *The chicken with the highest outdegree in an n-chicken*

*tournament  $G = (V, E)$  with  $n \geq 1$  is king.*

*Proof.* By contradiction. Let  $u$  be the node in the tournament graph with highest outdegree and

suppose that  $u$  is not a king. Let  $Y = \{v \mid u \rightarrow v \in E\}$  be the set of chickens that chicken  $u$

*pecks. Then  $\deg(u) = |Y|$ .*

*i.e.,*

Since  $u$  is not a king, there is a chicken  $x \notin Y$  (that is, that is not pecked by chicken  $u$ ) and

that is not pecked by any chicken in  $Y$ . Since for any pair of chickens, one pecks the other, this

means that  $x$  pecks  $u$  as well as every chicken in  $Y$ . This means that

$$\begin{aligned} &\text{outdegree}(x) \\ &\deg(x) = |Y| + 1 \\ &> \text{outdegree}(u), \\ &= \deg(u) + 1 \end{aligned}$$

$$> \deg(u)$$

But  $u$  was assumed to be the node with the highest degree in the tournament, so we have a

Redraft graphic: Missing graphic

Figure 6.7: A 5-chicken tournament in which every chicken is a king.

contradiction. Hence,  $u$  must be a king. ■

Theorem 6.2.2 means that if the player with the most victories is defeated by another player  $x$ , then at least he/she defeats some third player that defeats  $x$ . In this sense, the player with the most victories has some sort of bragging rights over every other player. Unfortunately, as Figure 6.6 illustrates, there can be many other players with such bragging rights, even some with fewer victories. Indeed, for some tournaments, it is possible that every player is a “king.” For example, consider the tournament illustrated in Figure 6.7.

### 6.3 Communication Networks

While reasoning about chickens pecking each other may be amusing (to mathematicians, at least), the use of directed graphs to model communication networks is very serious business. In the context of communication problems, vertices represent computers, processors, or switches, and edges represent wires, fiber, or other transmission lines through which data flows. For some communication networks, like the internet, the corresponding graph is enormous and largely chaotic. Highly structured networks, such as an array or butterfly, by contrast, find application in telephone switching systems and the communication hardware inside parallel computers.

X

### 6.3.1 Packet Routing

Whatever architecture is chosen, the goal of a communication network is to get data from *inputs*

to *outputs*. In this text, we will focus on a model in which the data to be communicated is in the

form of a *packet*. In practice, a packet would consist of a fixed amount of data, and a message

(such as a web page or a movie) would consist of many packets.

For simplicity, we will restrict our attention to the scenario where there is just one packet at

every input and where there is just one packet destined for each output. We will denote the

number of inputs and output by  $N$  and we will often assume that  $N$  is a power of two.

We will specify the desired destinations of the packets by a permutation<sup>1</sup> of  $\{0, 1, \dots, N - 1\}$

X

So a permutation,  $\pi$ , defines a *routing problem*: get a packet that starts at input  $i$  to output  $\pi(i)$

for  $0 \leq i < N$ . A *routing*,  $P$ , that *solves* a routing problem,  $\pi$ , is a set of paths from each input

to its specified output. That is,  $P$  is a set of paths,  $P_i$ , for  $i = 0, \dots, N - 1$ , where  $P_i$  goes from

X

<sup>1</sup>A permutation of a sequence is a reordering of the sequence.

X

input  $i$  to output  $\pi(i)$ .

Of course, the goal is to get all the packets to their destinations as quickly as possible using as little hardware as possible. The time needed to get the packages to their destinations depends on several factors, such as how many switches they need to go through and how many packets will need to cross the same wire. We will assume that only one packet can cross a wire at a time. The complexity of the hardware depends on factors such as the number of switches needed and the size of the switches.

Let's see how all this works with an example—routing packets on a complete binary tree.

### 6.3.2 The Complete Binary Tree

One of the simplest structured communications networks is a *complete binary tree*. A complete

binary tree with 4 inputs and 4 output is shown in Figure 6.8.

In this diagram and many that follow, the squares represent *terminals* (that is, the inputs and

OK  
as was

X

**Redraft graphic: Missing graphic**

Figure 6.8: A 4-input, 4-output complete binary tree. The squares represent terminals (in ~~put and~~

~~put~~/output registers) and the circles represent switches. Directed edges represent communication

channels in the network through which data packets can move. The unique path from input 1 to

output 3 is shown in bold.

outputs), and the circles represent *switches*, which direct packets through the network. A switch

receives packets on incoming edges and relays them forward along the outgoing edges. Thus,

you can imagine a data packet hopping through the network from an input terminal, through a

sequence of switches joined by directed edges, to an output terminal.

Recall that there is a unique simple path between every pair of vertices in a tree. So the natural

way to route a packet of data from an input terminal to an output in the complete binary tree is

along the corresponding directed path. For example, the route of a packet traveling from input 1

to output 3 is shown in bold in Figure 6.8.

### 6.3.3 Network Diameter

The delay between the time that a packet arrives at an input and arrives at its designated output ~~is known as~~ <sup>referred to as</sup>

~~is known as~~ the *latency* and it is a critical issue in communication networks. If congestion is

not a factor, then this delay is generally proportional to the length of the path a packet follows.

Assuming it takes one time unit to travel across a wire, and that there are no additional delays at

switches, the delay of a packet will be the number of wires it crosses going from input to output.<sup>2</sup>

Generally packets ~~are~~ <sup>a</sup> ~~in~~ <sup>15</sup> ~~in~~ <sup>using</sup> ~~by~~ <sup>possible</sup> ~~covered~~ <sup>path</sup> ~~possible~~ ~~path~~.

a shortest path routing, the worst ~~case~~ delay is the distance between the input and output that

<sup>2</sup>Latency can also be measured as the number of switches that a packet must pass through when traveling between

the most distant input and output, since switches usually have the biggest impact on network speed. For example, in the

complete binary tree example, the packet traveling from input 1 to output 3 crosses 5 switches, which is 1 less than the

number of edges traversed.

~~③ The distance between two nodes taken to be the length of the shortest path~~

~~X~~ the length of this shortest path is the distance between the input and output.

X

are farthest apart. This is called the *diameter* of the network. In other words, the diameter of a network<sup>3</sup> is the maximum length of any shortest path between an input and an output. For example, in the complete binary tree shown in Figure 6.8, the distance from input 1 to output 3 is six. No input and output are farther apart than this, so the diameter of this tree is also six.

More generally, the diameter of a complete binary tree with  $N$  inputs and outputs is  $2 \log N +$

2. (All logarithms in this lecture—and in most of computer science—are base 2.) This is quite good, because the logarithm function grows very slowly. We could connect up  $2^{20} = 1,048,576$

inputs and outputs using a complete binary tree and the worst input-output delay for any packet

would be this diameter, namely,  $2 \log(2^{20}) + 2 = 22$  42.

<sup>3</sup>The usual definition of *diameter* for a general graph (simple or directed) is the largest distance between *any* two vertices, but in the context of a communication network we're only interested in the distance between inputs and outputs, not between arbitrary pairs of vertices.

no italics

#### 6.3.4 Switch Size

One way to reduce the diameter of a network (and hence the latency needed to route packets)

is to use larger switches. For example, in the complete binary tree, most of the switches have

three incoming edges and three outgoing edges, which makes them  $3 \times 3$  switches. If we had

$4 \times 4$  switches, then we could construct a complete *ternary* tree with an even smaller diameter.

In principle, we could even connect up all the inputs and outputs via a single monster  $N \times N$

switch, as shown in Figure 6.9. In this case, the “network” would consist of a single switch and

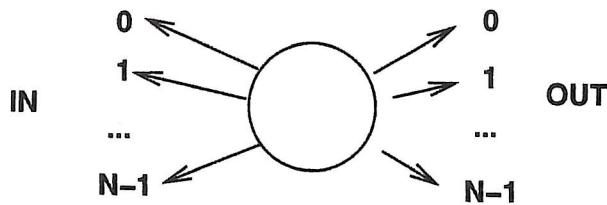
the latency would be 2.

This isn’t very productive, however, since we’ve just concealed the original network design

problem inside this abstract switch. Eventually, we’ll have to design the internals of the monster  
*Monster*  
^

switch using simpler components, and then we’re right back where we started. So the challenge

in designing a communication network is figuring out how to get the functionality of an  $N \times N$



**Redraft graphic: Fix labels**

Figure 6.9: A monster  $N \times N$  switch.

switch using fixed size, elementary devices, like  $3 \times 3$  switches.

### 6.3.5 Switch Count

Another goal in designing a communication network is to use as few switches as possible. The number of switches in a complete binary tree is  $1 + 2 + 4 + 8 + \dots + N = 2N - 1$ , since there is 1 switch at the top (the “root switch”), 2 below it, 4 below those, and so forth. This is nearly the best possible with  $3 \times 3$  switches, since at least one switch will be needed for each pair of inputs and outputs.

### 6.3.6 Congestion

The complete binary tree has a fatal drawback: the root switch is a bottleneck. At best, this

switch must handle an enormous amount of traffic: every packet traveling from the left side of

the network to the right or vice-versa. Passing all these packets through a single switch could

take a long time. At worst, if this switch fails, the network is broken into two equal-sized pieces.

# the traffic through the root depends on the routing problem.

For example, if the routing problem is given by the identity permutation,  $\text{id}(i) := i$ , then there

is an easy routing,  $P$ , that solves the problem: let  $P_i$  be the path from input  $i$  up through one

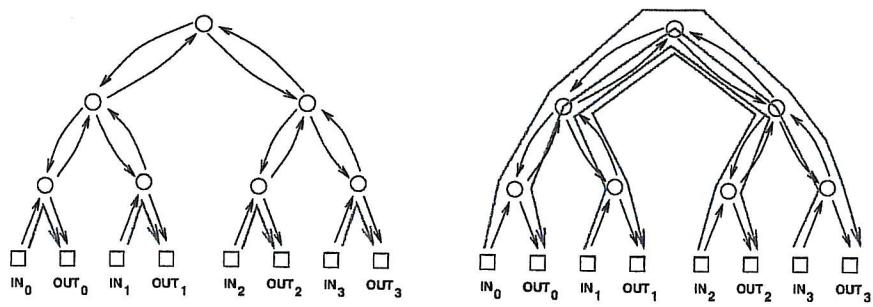
switch and back down to output  $i$ . On the other hand, if the problem was given by  $\pi(i) := (N -$

$P$   $P_i$   
1) -  $i$ , then in any solution,  $\mathcal{Q}$ , for  $\pi$ , each path  $\mathcal{Q}_i$  beginning at input  $i$  must eventually loop all

the way up through the root switch and then travel back down to output  $(N - 1) - i$ . These two

situations are illustrated in Figure 6.10.

We can distinguish between a “good” set of paths and a “bad” set based on congestion. The



**Redraft graphic: Add sublabels**

Figure 6.10: Paths for the routing problem given by  $\pi(i) = i$  (a) and  $\pi(i) = N - 1 - i$  (b). The root is a bottleneck in the latter scenario.

*congestion* of a routing,  $P$ , is equal to the largest number of paths in  $P$  that pass through a single switch. For example, the congestion of the routing in Figure 6.10(a) is 1, since at most 1 path passes through each switch. However, the congestion of the routing in Figure 6.10(b) is 4, since 4 paths pass through the root switch (and the two switches directly below the root). Generally, lower congestion is better since packets can be delayed at an overloaded switch.

By extending the notion of congestion to networks, we can also distinguish between “good” and “bad” networks with respect to bottleneck problems. For each routing problem,  $\pi$ , for the network, we assume a routing is chosen that optimizes congestion, that is, that has the minimum congestion among all routings that solve  $\pi$ . Then the largest congestion that will ever be suffered by a switch will be the maximum congestion among these optimal routings. This “maxi-min” congestion is called the *congestion of the network*.

You may find it helpful to think about max congestion in terms of a value game. You design your spiffy, new communication network; this defines the game. Your opponent makes the first

X

move in the game: she inspects your network and specifies a permutation routing problem that will strain your network. You move second: given her specification, you choose the precise paths that the packets should take through your network; you're trying to avoid overloading any one switch. Then her next move is to pick a switch with as large as possible a number of packets passing through it; this number is her score in the competition. The max congestion of your network is the largest score she can ensure; in other words, it is precisely the max-value of this game.

For example, if your enemy were trying to defeat the complete binary tree, she would choose a permutation like  $\pi(i) = (N - 1) - i$ . Then for *every* packet  $i$ , you would be forced to select a path  $P_{i,\pi(i)}$  passing through the root switch. Then, *you* enemy would choose the root switch and achieve a score of  $N$ . In other words, the max congestion of the complete binary tree is  $N$ —which is horrible!

We have summarized the results of our analysis of the complete binary tree in Figure 6.11.

network	diameter	switch size	# switches	congestion
complete binary tree	$2 \log N + 2$	$3 \times 3$	$2N - 1$	$N$
2-D array	$2N$	$2 \times 2$	$N^2$	2

Figure 6.11: A summary of the attributes of the complete binary tree.

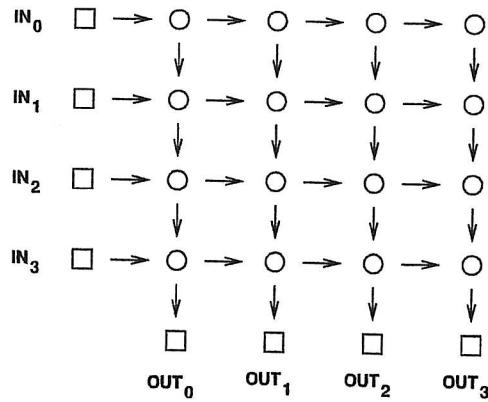
Overall, the complete binary tree does well in every category except the last—congestion, and that is a killer in practice. Next, we will look at a network that solves the congestion problem, but at a very high cost.

### 6.3.7

### 6.4 The 2-d Array ← subsection

An illustration of the  $N \times N$  2-d array (also known as the *grid* or *crossbar*) is shown in Figure 6.12 for the case when  $N = 4$ .

The diameter of the  $4 \times 4$  2-d array is 8, which is the number of edges between input 0 and

Figure 6.12: A  $4 \times 4$  2-dimensional array.

output 3. More generally, the diameter of a 2-d array with  $N$  inputs and outputs is  $2N$ , which is much worse than the diameter of the complete binary tree ( $2 \log N + 2$ ). On the other hand, replacing a complete binary tree with a 2-d array almost eliminates congestion.

**Theorem 6.4.1.** *The congestion of an  $N$ -input 2-d array is 2.*

*Proof.* First, we show that the congestion is at most 2. Let  $\pi$  be any permutation. Define a solution,  $P$ , for  $\pi$  to be the set of paths,  $P_i$ , where  $P_i$  goes to the right from input  $i$  to column

$\pi(i)$  and then goes down to output  $\pi(i)$ . In this solution, the switch in row  $i$  and column  $j$

~~encounters~~

~~transmits~~ at most two packets: the packet originating at input  $i$  and the packet destined for output  $j$ .

Next, we show that the congestion is at least 2. This follows because in any routing problem,

$\pi$ , where  $\pi(0) = 0$  and  $\pi(N - 1) = N - 1$ , two packets must pass through the lower left

switch. ■

The characteristics of the 2-d array are recorded in Figure 6.13. The crucial entry in this table

is the number of switches, which is  $N^2$ . This is a major defect of the 2-d array; a network with

$N = 1000$  inputs would require a *million*  $2 \times 2$  switches! Still, for applications where  $N$  is

small, the simplicity and low congestion of the array make it an attractive choice.

network	diameter	switch size	# switches	congestion
complete binary tree	$2 \log N + 2$	$3 \times 3$	$2N - 1$	$N$
2-D array	$2N$	$2 \times 2$	$N^2$	2

Figure 6.13: Comparing the  $N$ -input 2-d array to the  $N$ -input complete binary tree.

Redraft graphic: Missing graphic

Figure 6.14: An 8-input/output butterfly.

## 6.3.8

### 6.4.1 The Butterfly

The Holy Grail of switching networks would combine the best properties of the complete binary

tree (low diameter, few switches) and the array (low congestion). The *butterfly* is a widely-used

compromise between the two. A butterfly network with  $N = 8$  inputs is shown in Figure 6.14.

The structure of the butterfly is certainly more complicated than that of the complete binary or

2-d array. Let's see how it is constructed.

All the terminals and switches in the network are in  $N$  rows. In particular, input  $i$  is at the left

end of row  $i$ , and output  $i$  is at the right end of row  $i$ . Now let's label the rows in *binary* so that

the label on row  $i$  is the binary number  $b_1 b_2 \dots b_{\log N}$  that represents the integer  $i$ .

*S*  
Between the inputs and output, there are  $\log(N) + 1$  levels of switches, numbered from 0

to  $\log N$ . Each level consists of a column of  $N$  switches, one per row. Thus, each switch in the

network is uniquely identified by a sequence  $(b_1, b_2, \dots, b_{\log N}, l)$ , where  $b_1 b_2 \dots b_{\log N}$  is the

switch's row in binary and  $l$  is the switch's level.

All that remains is to describe how the switches are connected up. The basic connection pattern

is expressed below in a compact notation:

$$(b_1, b_2, \dots, b_{l+1}, \dots, b_{\log N}, l) \begin{cases} \nearrow & (b_1, b_2, \dots, b_{l+1}, \dots, b_{\log N}, l+1) \\ \searrow & (b_1, b_2, \dots, \overline{b_{l+1}}, \dots, b_{\log N}, l+1) \end{cases}$$

This says that there are directed edges from switch  $(b_1, b_2, \dots, b_{\log N}, l)$  to two switches in the

next level. One edges leads to the switch in the *same* row, and the other edge leads to the switch

*the  $(l+1)$ st bit  $b_{l+1}$ .*

in the row obtained by *inverting bit  $l + 1$* . For example, referring back to the illustration of the

size  $N = 8$  butterfly, there is an edge from switch  $(0, 0, 0, 0)$  to switch  $(0, 0, 0, 1)$ , which is in

the same row, and to switch  $(1, 0, 0, 1)$ , which is in the row obtained by inverting bit  $l + 1 = 1$ .

The butterfly network has a recursive structure; specifically, a butterfly of size  $2N$  consists of

two butterflies of size  $N$ , which are shown in dashed boxes in Figure 6.15, and one additional  
*additional*

level of switches. Each switch in the new level has directed edges to a pair of corresponding

*each of*  
switches in the smaller butterflies; one example is dashed in the figure. *For example, see ↗.*

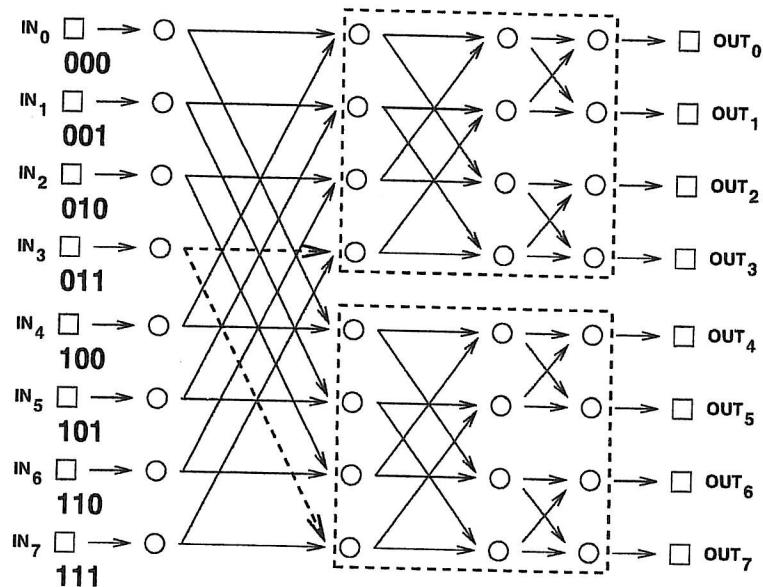
Despite the relatively complicated structure of the butterfly, there is a simple way to route

packets through its switches. In particular, suppose that we want to send a packet from input

$x_1 x_2 \dots x_{\log N}$  to output  $y_1 y_2 \dots y_{\log N}$ . (Here we are specifying the input and output numbers

*the first level*

in binary.) Roughly, the plan is to “correct” the first bit on *level 1*, correct the second bit on



Redraft graphic: Missing graphic

(shown in the dashed boxes).

Each switch  
the connection

Figure 6.15: An  $N$ -input butterfly contains two  $N/2$ -input butterflies. ~~the connection~~  
on the first level is ~~connected~~<sup>adjacent</sup> to a corresponding  
switch in each of the sub-butterflies. For  
example, we have used dashed lines to show  
the ~~the~~ edges for the ~~first level node~~  
node (0,1,1,0).

X

*The second level,*

~~level 2~~, and so forth. Thus, the sequence of switches visited by the packet is:

$$(x_1, x_2, x_3, \dots, x_{\log N}, 0) \rightarrow (y_1, x_2, x_3, \dots, x_{\log N}, 1)$$

$$\rightarrow (y_1, y_2, x_3, \dots, x_{\log N}, 2)$$

$$\rightarrow (y_1, y_2, y_3, \dots, x_{\log N}, 3)$$

$\rightarrow \dots$

$$\rightarrow (y_1, y_2, y_3, \dots, y_{\log N}, \log N)$$

In fact, this is the *only* path from the input to the output!

The congestion of the butterfly network is about  $\sqrt{N}$ . More precisely, the congestion is  $\sqrt{N}$

if  $N$  is an even power of 2 and  $\sqrt{N}/2$  if  $N$  is an odd power of 2. The task of proving this fact

has been left to the problem section. A comparison of the butterfly with the complete binary tree

and the 2-d array is provided in Figure 6.16. As you can see, the butterfly has lower congestion

than the complete binary tree. And it uses fewer switches and has lower diameter than the

*End* 1 The ~~permutation~~ routing problems that ~~result~~ in  $\sqrt{N}$  congestion do arise in practice, but for most routing problems, the congestion is much lower (around  $\log N$ ), which ~~is one reason why~~ ~~means that~~ the butterfly is ~~completely~~ useful in practice.

network	diameter	switch size	# switches	congestion
complete binary tree	$2 \log N + 2$	$3 \times 3$	$2N - 1$	$N$
2-D array	$2N$	$2 \times 2$	$N^2$	2
butterfly	$\log N + 2$	$2 \times 2$	$N(\log(N) + 1)$	$\sqrt{N}$ or $\sqrt{N/2}$

Figure 6.16: A comparison of the  $N$ -input butterfly with the  $N$ -input complete binary tree and the  $N$ -input 2-d array.

array. However, the butterfly does not capture the best qualities of each network, but rather is a compromise somewhere between the two. So our quest for the Holy Grail of routing networks goes on.

X

**6.3.9****6.5 Beneš Network** ← subsection

In the 1960's, a researcher at Bell Labs named Václav Beneš had a remarkable idea. He obtained

a marvelous communication network with congestion 1 by placing *two* butterflies back-to-back.

For example, the 8-input Beneš network is shown in Figure 6.17.

*roughly*

Putting two butterflies back-to-back *doubles* the number of switches and the diameter of a single butterfly, but it completely eliminates congestion problems! The proof of this fact relies on a

clever induction argument that we'll come to in a moment. Let's first see how the Beneš network

stacks up against the other networks we have been studying. As you can see in Figure 6.18, the

Beneš network has small size and diameter, and completely eliminates congestion.

*[This is a duck but true OK but it's a big deal.]* The Holy Grail of routing networks is in hand!

This is a duck but true OK but it's a big deal.

**Theorem 6.5.1.** *The congestion of the  $N$ -input Beneš network is 1 for any  $N$  that is a power*

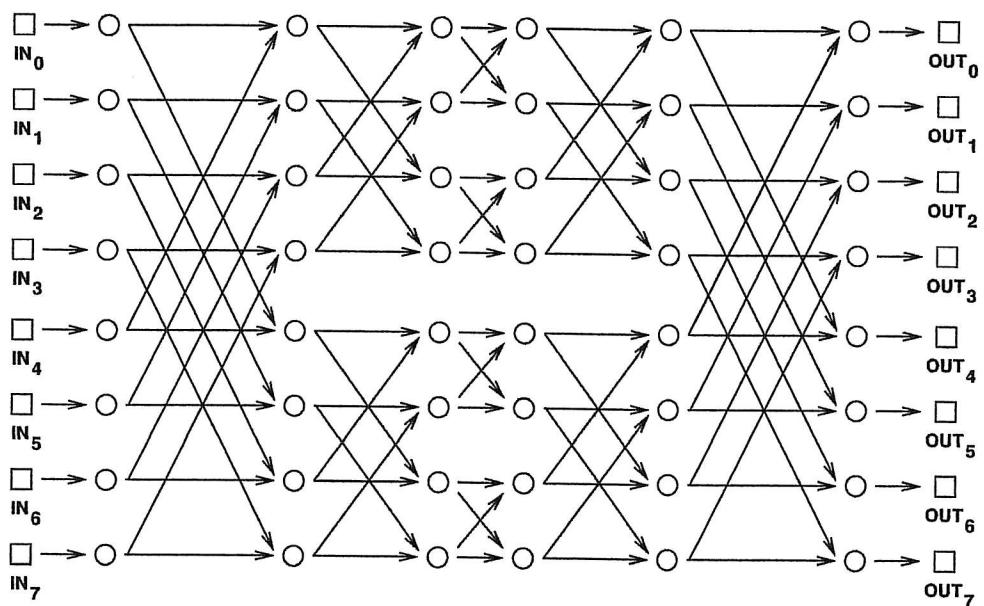


Figure 6.17: The 8-input Beneš network.

network	diameter	switch size	# switches	congestion
complete binary tree	$2 \log N + 2$	$3 \times 3$	$2N - 1$	$N$
2-D array	$2N$	$2 \times 2$	$N^2$	2
butterfly	$\log N + 2$	$2 \times 2$	$N(\log(N) + 1)$	$\sqrt{N}$ or $\sqrt{N/2}$
Beneš	$2 \log N + 1$	$2 \times 2$	$2N \log N$	1

Figure 6.18: A comparison of the  $N$ -input Beneš network with the  $N$ -input complete binary tree,

2-d array, and butterfly.

**Redraft graphic: Missing graphic**

Figure 6.19: The 2-input Beneš network.

of 2.

*Proof.* We use induction. Let  $P(a)$  be the proposition that the congestion of the  $2^a$ -input Beneš network is 1.

**Base case ( $a = 1$ ):** We must show that the congestion of the  $2^1$ -input Beneš network is 1. The network is shown in Figure 6.19.

There are only two possible permutation routing problems for a 2-input network. If  $\pi(0) = 0$  and  $\pi(1) = 1$ , then we can route both packets along the straight edges. On the other hand, if  $\pi(0) = 1$  and  $\pi(1) = 0$ , then we can route both packets along the diagonal edges. In both cases, a single packet passes through each switch.

**Inductive step:** We must show that  $P(a)$  implies  $P(a + 1)$  where  $a \geq 1$ . Thus, we assume

X

that the congestion of a  $2^a$ -input Beneš network is 1 in order to prove that the congestion of a  $2^{a+1}$ -input Beneš network is also 1.

### Digression

Time out! Let's work through an example, develop some intuition, and then complete the proof.

Notice that inside a Beneš network of size  $2N$  lurk two Beneš subnetworks of size  $N$ . This fol-

lows from our earlier observation that a butterfly of size  $2N$  contains two butterflies of size  $N$ . In

the Beneš network shown in Figure 6.20 with  $N = 8$  inputs and outputs, the two 4-input/output

shown

subnetworks are in dashed boxes.

↑

By the inductive assumption, the subnetworks can each route an arbitrary permutation with

congestion 1. So if we can guide packets safely through just the first and last levels, then we can

rely on induction for the rest! Let's see how this works in an example. Consider the following

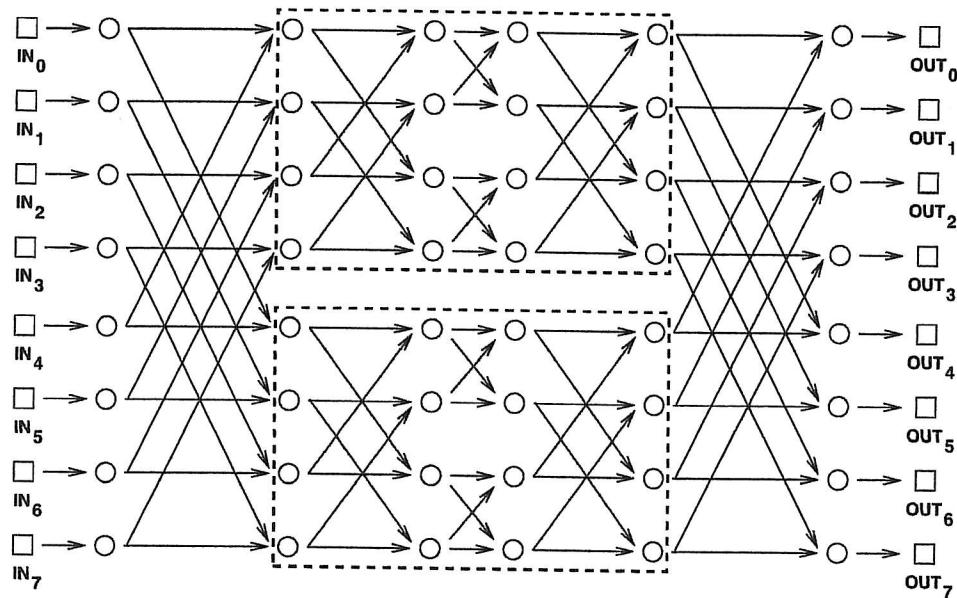


Figure 6.20: A  $2N$ -input Beneš network contains two  $N$ -input Beneš networks—shown here for

$$N = 4$$

X

permutation routing problem:

$$\pi(0) = 1 \quad \pi(4) = 3$$

$$\pi(1) = 5 \quad \pi(5) = 6$$

$$\pi(2) = 4 \quad \pi(6) = 0$$

$$\pi(3) = 7 \quad \pi(7) = 2$$

We can route each packet to its destination through either the upper subnetwork or the lower

subnetwork. However, the choice for one packet may constrain the choice for another. For

*the packets at inputs 0 and 4 both*

example, we can not route ~~both~~ *packet 0 and 4* through the same network since that

would cause two packets to collide at a single switch, resulting in congestion. So one packet

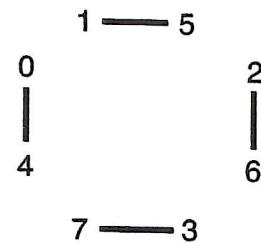
must go through the upper network and the other through the lower network. Similarly, *packets at inputs*

1 and 5, 2 and 6, and 3 and 7 must be routed through different networks. Let's record these

constraints in a graph. The vertices are the 8 packets. If two packets must pass through different

*(labeled according to their  
input position).*

K



The beginnings of a

Figure 6.21: A constraint graph for our packet routing problem. Adjacent packets cannot be  
routed using the same sub-Beneš network.

networks, then there is an edge between them. The resulting constraint graph is illustrated in

Figure 6.21. Notice that at most one edge is incident to each vertex.

The output side of the network imposes some further constraints. For example, the packet  
destined for output 0 (which is packet 6) and the packet destined for output 4 (which is packet  
2) can not both pass through the same network <sup>since</sup> that would require both packets to arrive from  
the same switch. Similarly, the packets destined for outputs 1 and 5, 2 and 6, and 3 and 7 must

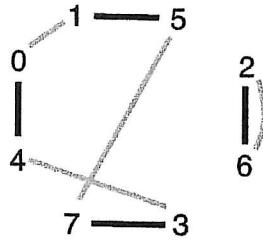


Figure 6.22: The updated constraint graph.

also pass through different switches. We can record these additional constraints in our constraint graph with gray edges, as is illustrated in Figure 6.22.

Notice that at most one new edge is incident to each vertex. The two lines drawn between vertices 2 and 6 reflect the two different reasons why these packets must be routed through different networks. However, we intend this to be a simple graph; the two lines still signify a single edge.

Now here's the key insight: *a 2-coloring of the graph corresponds to a solution to the routing*

problem. In particular, suppose that we could color each vertex either red or blue so that adjacent vertices are colored differently. Then all constraints are satisfied if we send the red packets through the upper network and the blue packets through the lower network.

The only remaining question is whether the constraint graph is 2-colorable. Fortunately, this is easy to verify:

*an undirected graph  $G$*

**Lemma 6.5.2.** *If the edges of ~~a graph~~ can be grouped into two sets such that every vertex is incident to at most 1 edge from each set, then the graph is 2-colorable.*

*Proof.* Since the two sets of edges may overlap, let's call an edge that is in both sets a *doubled edge*.

~~Consider a graph here.~~

~~We know from Theorem 5.9.1 that all we have to do is show that every closed walk has even length. There are two cases:~~

~~Case 1: The closed walk contains a doubled edge. No other edge can be incident to either of the~~

*Note that no*

endpoints of a doubled edge, since that endpoint would then be incident to two edges from

~~This means that doubled edges form the same set. So a closed walk traversing a doubled edge has nowhere to go but back and forth along the edge an even number of times.~~

~~These connected components with 2 nodes. Such connected components are easily colored with 2 colors and so we can henceforth ignore them, and focus on~~

~~Case 2: No edge on the closed walk is doubled. Since each vertex is incident to at most one edge~~

~~from each set, any path with no doubled edges must traverse successive edges that alternate~~

~~from one set to the other. In particular, a closed walk must traverse a path of alternating~~

~~edges that begins and ends with edges from different sets. This means the closed walk has~~

~~to be of even length.~~ ■

For example, a 2-coloring of the constraint graph in Figure 6.22 is shown in Figure 6.23. ↗

The solution to this graph-coloring problem provides a start on the packet routing problem.

We can complete the routing in the two smaller Beneš networks by induction. With this insight

~~The digression is over and~~

in hand, we can now complete the proof of Theorem 6.5.1. ↘

## INSERTA

The remaining nodes and edges, which form a simple graph.

DAVID - This is in the new <sup>sub</sup>section on "odd cycles and 2-colorability" in CHS

By Theorem XY, we know that ~~\_\_\_\_\_~~ if a simple graph has no odd cycles, then it is 2-colorable. So all we need to do is show that every cycle in ~~the constant graph~~ has even length. This is easy to do since any cycle on 6

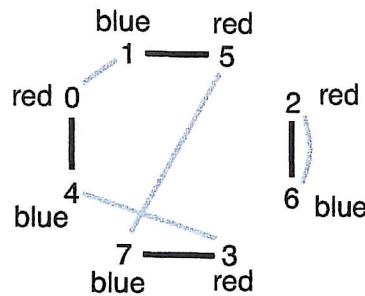


Figure 6.23: A 2-coloring of the constraint graph in Figure 6.22.

*Proof of Theorem 6.5.1 (continued).* Let  $\pi$  be an arbitrary permutation of  $\{0, 1, \dots, N - 1\}$ . Let

$G$  be the graph whose vertices are packet numbers  $0, 1, \dots, N - 1$  and whose edges come from

the union of these two sets:

$$E_1 := \{u-v \mid |u-v| = N/2\}, \text{ and}$$

$$E_2 := \{u-w \mid |\pi(u) - \pi(w)| = N/2\}.$$

big dash

Now any vertex,  $u$ , is incident to at most two edges: a unique edge  $u-v \in E_1$  and a unique edge  $u-w \in E_2$ . So according to Lemma 6.5.2, there is a 2-coloring for the vertices of  $G$ . Now

route packets of one color through the upper subnetwork and packets of the other color through the lower subnetwork. Since for each edge in  $E_1$ , one vertex goes to the upper subnetwork and the other to the lower subnetwork, there will not be any conflicts in the first level. Since for each edge in  $E_2$ , one vertex comes from the upper subnetwork and the other from the lower subnetwork, there will not be any conflicts in the last level. We can complete the routing within each subnetwork by the induction hypothesis  $P(n)$ . ■

---

*6.6* **6.4**  
Problems

