

Part II

~~Structures
Mathematical Data Types~~

— ~~INERT 5A goes here~~ —

INSERT 5A

how to

~~proof~~

Now that we understand what a ~~good~~ proof is and how to construct a good proof, we can turn our attention to the structures that are most useful in modelling problems that arise in computer science. The most important structure is a graph (aka network). Graphs ~~are~~ provide an excellent mechanism for modelling associations between pairs of objects (e.g., two exams that cannot be given at the same time, two people that like each other, or two subroutines that can be run independently). In Chapter 5, we study graphs that represent symmetric relationships, and in Chapter 6, we consider graphs ~~where~~ where the relationship is ~~one-way~~ (e.g., ~~one person of~~ ~~one set of~~ ~~operations where one needs to be done before another~~ ~~operations where one needs to be done before another~~ a situation where you can ^{go} from ~~get~~ ~~from~~ ~~x to y but not y to x~~ necessarily vice-

Versa.

In Chapter 7, we ~~examine~~ consider the more general notion of a relation and we examine important classes of relations such as ~~equivalence~~ and partially ordered sets. Partially ordered sets arise frequently in scheduling problems.

We conclude in Chapter 8 with a study of state machines. State machines can be used to model a variety of processes and ~~one~~ ^{are} fundamental tools in proving ^{that an algorithm} ~~the correctness~~ and termination ~~property~~ and that it produces the correct output.

CSA

Chapter 5

Graph Theory

~~Simple Graphs~~

*& then 5C go here —
5C go*

*— INSERTS 5B, ~~Notes from lesson~~
(5B is text from
P 358)*

~~Graphs in which edges are not directed are called simple graphs. They come up in~~

~~all sorts of applications, including scheduling, optimization, communications, and~~

~~the design and analysis of algorithms. Two Stanford students even used graph~~

~~theory to become multibillionaires!~~

~~But we'll start with an application designed to get your attention: we are going to make a professional inquiry into sexual behavior. Namely, we'll look at some data about who, on average, has more opposite-gender partners, men or women.~~

INSERT 5C

Graphs are ubiquitous in computer science ~~because~~ because they provide ~~a very complex way~~ a handy way to represent ~~a~~ relationship between pairs of objects. We simply ~~place an edge between~~ ~~simply put, we use and edge between~~ ~~a pair~~

The objects represent items of interest such as programs, people, cities, or web pages, and we ~~place an~~ place an edge between a pair of nodes if they are related in a certain way.

For example, an edge between a pair of people might indicate that they like (or, in ~~other~~ ^{alternate} scenarios, that they don't like) each other. An edge between a pair of courses might indicate that ~~you cannot take them~~

one needs to be taken before the other.

In this chapter, we will focus our attention on simple graphs where the relationship denoted by an edge is symmetric. Afterward, in Chapter 6, we consider the situation where the edge \Rightarrow denotes a one-way relationship (such e.g., where one web page points to the other.)

~~course needs to be taken before the other).~~

1 ~~Analyzing the properties of such a graph~~

Two Stanford students analyzed ~~such~~ such a graph to become multi-billionaires. ~~so they~~ So, pay ~~pay~~ attention to graph theory, and who knows what ~~it~~ might happen!

This is where it goes to page 371

Sexual demographics have been the subject of many studies. In one of the largest, researchers from the University of Chicago interviewed a random sample of 2500 ~~people~~ Americans over several years to try to get an answer to this question. Their study, published in 1994, and entitled *The Social Organization of Sexuality* found that on average men have 74% more opposite-gender partners than women.

Other studies have found that the disparity is even larger. In particular, ABC News claimed that the average man has 20 partners over his lifetime, and the average woman has 6, for a percentage disparity of 233%. The ABC News study, aired on Primetime Live in 2004, purported to be one of the most scientific ever done, with only a 2.5% margin of error. It was called "American Sex Survey: A peek

between the sheets," which raises some question about the seriousness of their reporting.

~~EDITING NOTE:~~ The promotion for the study is even better:

"A ground breaking ABC News "Primetime Live" survey finds a range of eye-popping sexual activities, fantasies and attitudes in this country,

confirming some conventional wisdom, exploding some myths – and

" venturing where few scientific surveys have gone before.

Probably that last part about going where few scientific surveys have gone before

is pretty accurate!

Yet again, in August, 2007, the N.Y. Times reported on a study by the National

Center for Health Statistics of the U.S. government showing that men had seven

partners while women had four. Anyway, whose numbers do you think are more

accurate, the University of Chicago, ABC News, or the National Center —don't

answer; this is a setup question like "When did you stop beating your wife?" Using

we will now

a little graph theory, we'll explain why none of these findings can be anywhere

near the truth.

Definitions

9.1 Degrees & Isomorphism

9.1.1 Definition of Simple Graph 9.1.1 Simple Graphs

Informally, a graph is a bunch of dots with lines connecting some of them. Here

shown in Figure AA. The dots are called nodes (or vertices) and the lines are called edges.

An example is

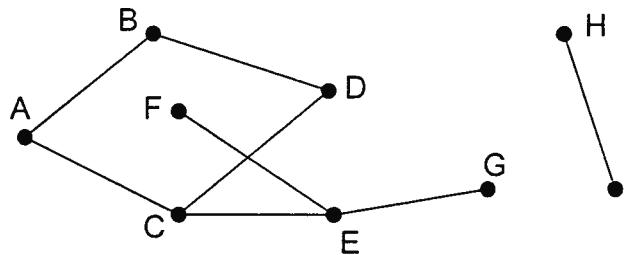


Figure AA: An example of a graph with 9 nodes and 8 edges.

For many mathematical purposes, we don't really care how the points and lines

are laid out — only which points are connected by lines. The definition of simple

graph aims to capture just this connection data.

(a.k.a. nodes)
for nodes

Definition 9.1.1. A simple graph, G , consists of a nonempty set, V , called the vertices

of G , and a collection, E , of two-element subsets of V . The members of E are called

the edges of G , and we write $G = (V, E)$.

Figure AA

The vertices correspond to the dots in Figure AA and the edges correspond to

1 we will use the term vertex and node interchangeably.

9.1. DEGREES & ISOMORPHISM

359

The graph in Figure 4A is written mathematically
the lines. For example, the connection data given in the diagram above can also be

given by listing the vertices and edges according to the official definition of simple
as $G = (V, E)$ where:

~~graph:~~

$$V = \{A, B, C, D, E, F, G, H, I\}$$

$$E = \{\{A, B\}, \{A, C\}, \{B, D\}, \{C, D\}, \{C, E\}, \{E, F\}, \{E, G\}, \{H, I\}\}.$$

In this case, the graph $G = (V, E)$ has 9 nodes and 8 edges.

It will be helpful to use the notation $A-B$ for the edge $\{A, B\}$. Note that $A-B$
~~often~~

and $B-A$ are different descriptions of the same edge, since sets are unordered.

So the definition of simple graphs is the same as for directed graphs, except
that instead of a directed edge $v \rightarrow w$ which starts at vertex v and ends at vertex
 w , a simple graph only has an undirected edge, $v-w$, that connects v and w .

Definition 9.1.2. Two vertices in a simple graph are said to be *adjacent* if they are

joined by an edge, and an edge is said to be *incident* to the vertices it joins. The

number of edges incident to a vertex is called the *degree* of the vertex v equivalently,

the degree of a vertex is equals the number of vertices adjacent to it.

For example, in the simple graph above, A is adjacent to B and B is adjacent to

D , and the edge $A-C$ is incident to vertices A and C . Vertex H has degree 1, D

$$\cancel{\text{deg}(E)=3}.$$

has degree 2, and E has degree 3. It is possible for a vertex to have degree 0, in which case it is not adjacent to any other vertex. A simple graph does not need ~~any other edges~~ vertices.

~~Graph Synonyms~~ to have any edges at all (in which case, the degree of every vertex is zero and $|E|=0$), but it

A synonym for "vertices" is "nodes," and we'll use these words interchangeably.

does ~~not~~ need to have at least one vertex.

Simple graphs are sometimes called networks, edges are sometimes called arcs. We

(i.e., $|V| \geq 1$).

Mention this as a "heads up" in case you look at other graph theory literature; we

won't use these words.

Some technical consequences of Definition 9.1.1 are worth noting right from the

start:

Note that simple

Simple graphs do not have self-loops (a,a is not an undirected edge between a and a).

An undirected edge is defined to be a set of two vertices. In addition, there

2. There is at most one edge between two vertices of a simple graph.

That is because E is a set. Lastly, simple graphs do not contain

3. Simple graphs have at least one vertex, though they might not have any

directed edges (i.e., edges of the form (a,b) instead of $\{a,b\}$).

2 Recall that the notation $|E|$ denotes the cardinality of the set E (i.e., the number of elements in E).

In other words, a simple graph does not contain multiple edges or multiple loops.

#

There's no harm in relaxing these conditions, and some authors do, but we don't

need self-loops, multiple edges between the same two vertices, or graphs with no

vertices, and it's simpler not to have them around. *we will consider graphs with directed edges (called directed graphs or digraphs) at length in Chapter 10.* For the rest of this chapter we'll only be considering simple graphs, so we'll just call them "graphs" from now on. *(in this chapter)*

This is from section 9.1
and goes to page 371

9.1.2 Sex in America

Let's model the question of heterosexual partners in graph theoretic terms. To do

this, we'll let G be the graph whose vertices, V , are all the people in America.

Then we split V into two separate subsets: M , which contains all the males, and

F , which contains all the females.¹ We'll put an edge between a male and a female

iff they have been sexual partners. This graph is pictured in Figure 9.1 with males

on the left and females on the right.

G

Actually, this is a pretty hard graph to figure out, let alone draw. The graph

is enormous: the US population is about 300 million, so $|V| \approx 300M$. Of these,

In the United States,

¹For simplicity, we'll ignore the possibility of someone being both, or neither, a man and a woman.

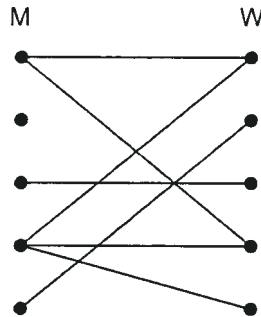


Figure 9.1: The sex partners graph

of the population is
 approximately 50.8% are female and 49.2% are male, so $|M| \approx 147.6M$, and $|F| \approx$

152.4M. And we don't even have trustworthy estimates of how many edges there

are, let alone exactly which couples are adjacent. But it turns out that we don't

to debunk the sex surveys

need to know any of this—we just need to figure out the relationship between

the average number of partners per male and partners per female. To do this,

and one F vertex

we note that every edge is incident to exactly one M vertex (remember, we're only

considering male-female relationships); so the sum of the degrees of the M vertices

and

equals the number of edges. For the same reason, the sum of the degrees of the F

of course

vertices equals the number of edges. So these sums are equal:

$$\sum_{x \in M} \deg(x) = \sum_{y \in F} \deg(y).$$

~~← If~~

~~Now suppose we divide both sides of this equation by the product of the sizes of~~

~~the two sets, $|M| \cdot |F|$; we obtain:~~

$$\left(\frac{\sum_{x \in M} \deg(x)}{|M|} \right) \cdot \frac{1}{|F|} = \left(\frac{\sum_{y \in F} \deg(y)}{|F|} \right) \cdot \frac{1}{|M|} \quad (\text{Eqn 5A})$$

~~— INSERT 5* goes here
open the size terms in Equation 5A~~

~~The terms above in parentheses are the average degree of an M vertex and the average~~

~~this means Hence, Equation 5A
degree of a F vertex. So we know:~~

~~implies that the average degree of an M vertex~~

~~is $\frac{|F|}{|M|}$ times the average degree of an F vertex.~~

~~In other words, we've proved that the average number of female partners of~~

~~males in the population compared to the average number of males per female is~~

~~determined solely by the relative number of males and females in the population.~~

~~From~~

~~we know~~

~~Now the Census Bureau reports that there are slightly more females than males~~

~~in America; in particular $|F|/|M|$ is about 1.035. So we know that on average,~~

~~males have 3.5% more opposite-gender partners than females. ~~about this tells us~~~~

~~of course, ~~it implies~~~~

~~This statistic really says~~

INSERT 50X

56⁻⁴

Notice that

$$\frac{\sum_{x \in M} \deg(x)}{|M|}$$

is ~~not~~ simply the average degree of a ~~node~~ node in M . This is ~~simply~~ the ~~average~~ average number of opposite-gender partners for a male in America. Similarly

$$\frac{\sum_{y \in F} \deg(y)}{|F|}$$

is the average degree of a node in F , which is the average number of opposite-gender partners for a female in America. Hence, Equation 5A implies that ~~the average male~~ ~~on average, an American male~~ has $|F|/|M|$ times as many opposite-gender partners as the average American female.

Remarkably, promiscuity is completely irrelevant in this analysis. That is because, the ratio of the average number of partners is completely determined by

364

CHAPTER 9. SIMPLE GRAPHS

the

~~that is totally irrelevant~~

nothing about any sex's promiscuity or selectivity. Rather, it just has to do with the

relative number of males and females. Collectively, males and females have the

same number of opposite gender partners, since it takes one of each set for every

partnership, but there are fewer males, so they have a higher ratio. This means

that the University of Chicago, ABC, and the Federal ~~60~~ government studies are way

off. After a huge effort, they gave a totally wrong answer.

There's no definite explanation for why such surveys are consistently wrong.

One hypothesis is that males exaggerate their number of partners —or maybe fe-

males downplay theirs —but these explanations are speculative. Interestingly, the

principal author of the National Center for Health Statistics study reported that

she knew the results had to be wrong, but that was the data collected, and her job

was to report it.

The same underlying issue has led to serious misinterpretations of other survey

data. For example, a couple of years ago, the Boston Globe ran a story on a survey

of the study habits of students on Boston area campuses. Their survey showed

that on average, minority students tended to study with non-minority students more than the other way around. They went on at great length to explain why this "remarkable phenomenon" might be true. But it's not remarkable at all —using our graph theory formulation, we can see that all it says is that there are fewer minority students than non-minority students, which is, of course what "minority" means.

The
~~9.1.3~~ Handshaking Lemma

↙ Subsubsection (note)

The previous argument hinged on the connection between a sum of degrees and the number edges. There is a simple connection between these in any graph:

Lemma 9.1.3. *The sum of the degrees of the vertices in a graph equals twice the number of edges.*

Proof. Every edge contributes two to the sum of the degrees, one for each of its endpoints. ■

Lemma 9.1.3 is sometimes called the *Handshake Lemma*: if we total up the number of people each person at a party shakes hands with, the total will be twice the number of handshakes that occurred.

9.1.2

9.1.4 Some Common Graphs

Some graphs come up so frequently that they have names. The *complete graph* on n

vertices, also called K_n , has an edge between every two vertices. ~~Here is K_5 :~~ shown in Figure A.

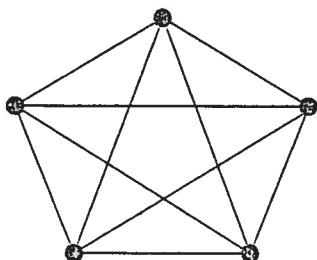


Figure A: The complete graph on 5 nodes, K_5 .

The *empty graph* has no edges at all. Here is the empty graph on 5 vertices:



~~An example~~ For example, the empty graph on 5 vertices with empty graph on 5 nodes is shown in Figure B.

Figure B: The empty graph with 5 nodes.

The n -node graph containing $n-1$ edges in sequence ~~in path form~~ is known as the line graph L_n . More formally $L_n = \{V, E\}$ where $V = \{v_1, v_2, \dots, v_n\}$ and

9.1. DEGREES & ISOMORPHISM $E \subseteq \{\{v_1, v_2\}, \{v_2, v_3\}, \dots, \{v_{n-1}, v_n\}\}$, For example, L_5 is displayed in Figure C. Another 5 vertex graph is C_5 , the line graph of length four.

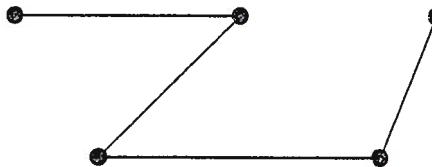


Figure C. The 5-node line graph L_5 .

~~The n-node graph~~

And here is C_5 , a simple cycle with 5 vertices.

If we add the edge $\{v_n, v_1\}$ to the line graph L_n ,

we get a ~~simplecycle~~

~~the graph consisting of a simple cycle.~~

For example, C_5 is illustrated in Figure D

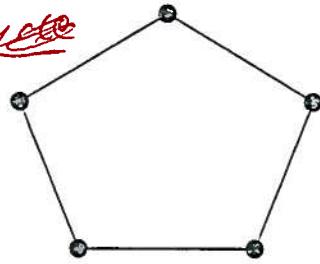


Figure D. The 5-node cyclograph C_5 .

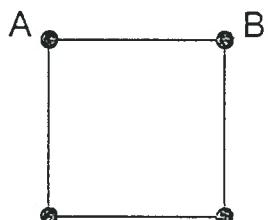
9.1.3

✓ 9.1.5 Isomorphism

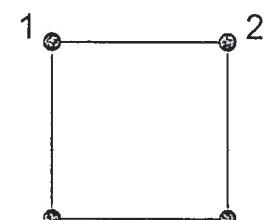
Two graphs that look the same might actually be different in a formal sense. For

in Figure E

example, the two graphs below are both simple cycles with 4 vertices, but



(a)



(b)

Figure E : Two graphs that are isomorphic to C_4 .

~~←~~ One graph has vertex set $\{A, B, C, D\}$ while the other has vertex set $\{1, 2, 3, 4\}$.

~~strictly speaking, these~~

~~if so, then the graphs are different mathematical objects, strictly speaking. But this~~

~~since~~

is a frustrating distinction: the graphs look the same!

~~but thus~~

through the

Fortunately, we can neatly capture the idea of "looks the same" by adapting

~~the notion of graph isomorphism.~~

~~Definition 6.2.1 of isomorphism of digraphs to handle simple graphs.~~

If $G_1 = (V_1, E_1)$ and $G_2 = (V_2, E_2)$ are two

~~Definition 9.1.4. If G is a graph with vertices, V , and edges, E , and likewise for graphs, then we say that~~

~~G_1 , then G_1 is isomorphic to G_2 iff there exists a bijection, $f : V_1 \rightarrow V_2$, such that~~

~~function~~
bijection¹

for every pair of vertices $u, v \in V_1$:

$$u - v \in E_1 \text{ iff } f(u) - f(v) \in E_2.$$

The function f is called an *isomorphism* between G_1 and G_2 .

EDITING NOTE: Graphs G_1 and G_2 are isomorphic if there exists a bijection be-

tween the vertices in G_1 and the vertices in G_2 such that there is an edge between

two vertices in G_1 if and only if there is an edge between the two corresponding

vertices in G_2 .

■

1 A bijection $f : V_1 \rightarrow V_2$ is a function that associates every node in V_1 with a unique node in V_2 and vice-versa. We will study bijections more deeply in Part III.

In other words, two graphs are isomorphic if ~~and~~ they are the same up to a relabeling of their vertices.

9.1. DEGREES & ISOMORPHISM

369

For example, here is an isomorphism between vertices in the two graphs ~~above~~ shown in Figure G:

A corresponds to 1
D corresponds to 4

B corresponds to 2
C corresponds to 3.

You can check that there is an edge between two vertices in the graph on the left if

and only if there is an edge between the two corresponding vertices in the graph

on the right.

Two isomorphic graphs may be drawn very differently. For example, ~~here are~~ we have shown

two different ways of drawing C_5 . ~~in Figure F.~~

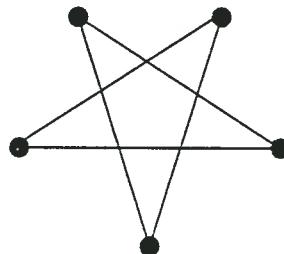
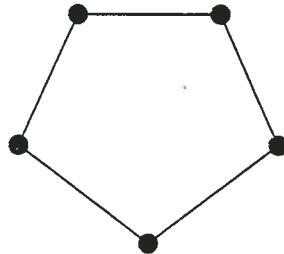


Figure F : Two ways of drawing C_5 .

Isomorphism preserves the connection properties of a graph, abstracting out

what the vertices are called, what they are made out of, or where they appear in a

drawing of the graph. More precisely, a property of a graph is said to be *preserved*

under isomorphism if whenever G has that property, every graph isomorphic to G

also has that property. For example, since an isomorphism is a bijection between

~~sets of vertices~~, isomorphic graphs must have the same number of vertices. What's

more, if f is a graph isomorphism that maps a vertex, v , of one graph to the ver-

tex, $f(v)$, of an isomorphic graph, then by definition of isomorphism, every vertex

adjacent to v in the first graph will be mapped by f to a vertex adjacent to $f(v)$

in the isomorphic graph. That is, v and $f(v)$ will have the same degree. So if one

graph has a vertex of degree 4 and another does not, then they can't be isomorphic.

In fact, they can't be isomorphic if the number of degree 4 vertices in each of the

graphs is not the same.

Looking for preserved properties can make it easy to determine that two graphs

are not isomorphic, or to actually find an isomorphism between them if there is

one. In practice, it's frequently easy to decide whether two graphs are isomorphic.

However, no one has yet found a *general* procedure for determining whether two

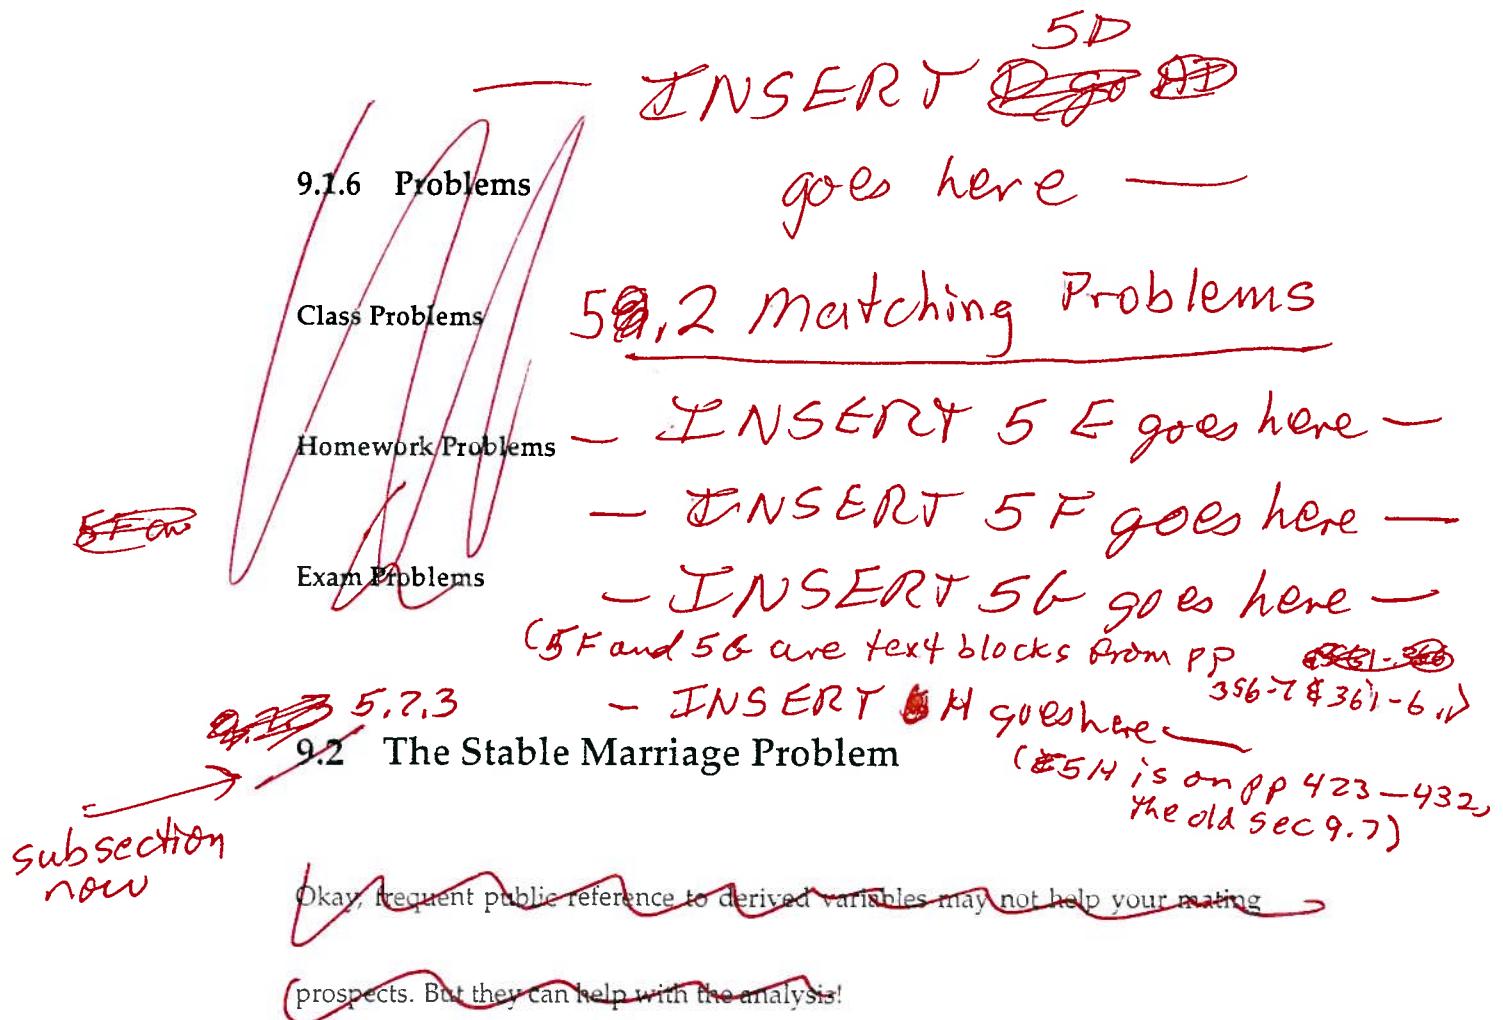
graphs are isomorphic that is *guaranteed* to run ~~much faster than an exhaustive~~

~~that is bounded by a power of the number~~
~~(and exhausting) search through all possible bijections between their vertices.~~

→ in polynomial time¹ in $|V|$.

1 i.e., in an amount of time that is upper bounded by $|V|^c$ where c is a ~~fixed~~ number independent of $|V|$.

Having an efficient procedure to detect isomorphic graphs would, for example, make it easy to search for a particular molecule in a database given the molecular bonds. On the other hand, knowing there is no such efficient procedure would also be valuable: secure protocols for encryption and remote authentication can be built on the hypothesis that graph isomorphism is computationally exhausting.



5G.1.4 Subgraphs

A graph $G_1 = (V_1, E_1)$

Definition: A graph $G_1 = (V_1, E_1)$ is said to be a subgraph of a graph $G_2 = (V_2, E_2)$ if $V_1 \subseteq V_2$ and $E_1 \subseteq E_2$.

For example, the empty graph on n nodes is a subgraph of L_n , L_n is a subgraph of C_n and C_n is a subgraph of K_n . Also, the graph $G_1 = (V, E)$ where

$$V = \{A, B, C\} \text{ and}$$

$$E = \{\{A, B\}, \{B, C\}\}$$

is a subgraph of the graph in Figure AA. On the other hand, the graph any graph containing an edge $\{G, H\}$ would not be a subgraph of the graph in Figure AA because the graph in Figure AA does not contain ~~this edge~~ ^{this edge}.

Note that since ~~the~~ a subgraph is itself a graph, the endpoints of any edge in a subgraph must also be in the subgraph. In other words if $G' = (V', E')$ is a subgraph of some graph G , and $\{v_i, v_j\} \in E'$, then it must be the case that $v_i \in V'$ and $v_j \in V'$.

5.1.5 Weighted Graphs

sometimes, we will use edges to denote a connection between a pair of nodes ~~where~~ where the connection has a capacity or weight. For example, ~~we might~~ ~~Internet fibers~~ ~~we might want~~ ~~connections~~, ~~pass of cities~~ ~~might~~ ~~be interested in~~; the capacity ~~two cities~~, ~~we could have different capacities~~, ~~resistors~~ ~~the resistance of a wire between two terminals~~, ~~the resistance of a wire between a pair of terminals~~, ~~right have different resistances~~, ~~the tension of a~~

spring connecting a pair of devices to a dynamical system, ~~right have different tensions~~ ~~of a~~ bonds between a pair of atoms. In a ~~a~~ molecule, ~~right~~ ~~or the distance of a highway between a pair of cities might have different strengths~~ ~~and so on~~.

In such cases, it is useful to represent the

~~system with a~~ weighted graph, ~~aka, a~~ weighted graph). A weighted graph is the same

as a simple graph except that we associate a real number (i.e., the weight) with each edge in the graph. Mathematically speaking, a weighted graph consists of a graph $G = (V, E)$ and a weight function $w: E \rightarrow \mathbb{R}$. For example, Figure 54 shows a weighted graph where the weight of edge A-B is 5.

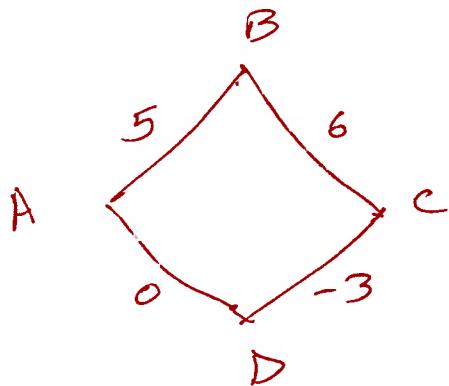


Figure 5Y: A 4-node weighted graph where edge $A \rightarrow B$ has weight 5.

5.1.6 Adjacency Matrices

There are many ways to represent a graph. A ~~for example~~, you can draw it ~~as shown in Figure 5Y~~, and you can represent it ~~as~~ with sets ~~as~~ (as in $G = (V, E)$). Another common representation is with an adjacency matrix.

We have already seen two ways: you

Definition 5AB: Given an n -node ~~graph~~ $G = (V, E)$, the adjacency matrix ~~of G~~ is

where $V = \{v_1, v_2, \dots, v_n\}$,

for G is the $n \times n$ matrix ~~$A_G = \{a_{ij}\}$~~ where

$$a_{ij} = \begin{cases} \cancel{1} & \cancel{\text{if } \{v_i, v_j\} \in E} \\ 1 & \text{if } \{v_i, v_j\} \in E \\ 0 & \text{otherwise} \end{cases}.$$

with edge weight given by $w: E \rightarrow \mathbb{R}$

If G is a weighted graph, then the adjacency matrix for G is ~~$A_G = \{a_{ij}\}$~~ where

$$a_{ij} = \begin{cases} f(\{v_i, v_j\}) & \text{if } \{v_i, v_j\} \in E \\ 0 & \text{otherwise} \end{cases}.$$

Figure 52 displays

For example, the adjacency matrices for the graphs shown in Figures E(a) and 5Y where $v_1 = A, v_2 = B, v_3 = C$ and $v_4 = D$.

	$\left(\begin{array}{cccc} 0 & 1 & 0 & 0 \\ 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \\ 1 & 0 & 0 & 0 \end{array} \right)$	$\left(\begin{array}{cccc} 0 & 5 & 0 & 0 \\ 5 & 0 & 6 & 0 \\ 0 & 6 & 0 & 3 \\ 0 & 0 & 3 & 0 \end{array} \right)$
(a)	(b)	

Figure 52 : Examples of the adjacency matrices. (a)

Shows the adjacency matrix for the graph in Figure E(a) and (b) shows the adjacency matrix for the weighted graph in Figure 5Y. In each case, we use the association set $v_1 = A, v_2 = B, v_3 = C$ and $v_4 = D$ to construct the matrix.

we begin our study of graph theory by considering the scenario where the nodes in a graph represent people and edges represent a relationship between pairs of people such as "likes", "marries", and so on. ~~Although~~ Now, ~~we will spend talk about the graph~~ ~~examples will generally be in terms of~~ ~~you may be wondering what~~ marriage has to do with computer science, ~~but~~, turns out that the techniques we will develop ~~will~~ apply to much more general scenarios ~~that~~ where instead of matching men to women, we need to match packets to paths in a network ~~or~~ applicants to jobs, or internet traffic to web servers.

In our first example, we will show how graph theory can be used to debunk ~~some~~ an urban legend about sexual ~~practices~~ practices in America. Yes, you read correctly. So, fasten your seat belt — ~~who knew that math~~ ~~to get at more interesting~~ ~~is~~ might actually be interesting!

INSERTED

5E-2

Our focus

9.2.1 Sex in America

On average, who has more opposite-gender partners: men or women?

9.2.1 The Problem

(Subsubsection (no #))

We next consider a version of the bipartite matching problem where there are an equal number of men and women, and where each person has preferences about who they would like to marry. In fact, we

~~Suppose there are a bunch of boys and an equal number of girls that we want to marry off. Each boy has his personal preferences about the girls—in fact, we~~

~~that each man women~~

assume he has a complete list of all the ~~girls~~ ranked according to his preferences,

~~A~~ ~~woman~~ ~~men~~.

with no ties. Likewise, each ~~girl~~ has a ranked list of all of the ~~boys~~.

The preferences don't have to be symmetric. That is, Jennifer might like Brad

everyone!

best, but Brad doesn't necessarily like Jennifer best. The goal is to marry off ~~boys~~

~~and girls: every boy must marry exactly one girl and vice-versa—no polygamy. Another word, we are looking for a perfect matching.~~

~~In mathematical terms, we want the mapping from boys to their wives to be a bijection or perfect matching. We'll just call this a "matching," for short.~~

~~There is no ~~couple~~ pair of people that ~~both~~ prefer each other to their spouses.~~

~~Here's the difficulty: suppose every boy likes Angelina best, and every girl likes~~

For example,

~~man~~

~~woman~~

Brad best, but Brad and Angelina are married to other people, say Jennifer and

Billy Bob. Now Brad and Angelina prefer each other to their spouses, which puts their

marriages at risk: pretty soon, they're likely to start spending late nights ~~studying~~

together

working on problem sets!

~~sessions together all the time~~

~~unfortunate~~ This situation is illustrated in ~~the following diagram~~ where the digits "1" and "2" near a ~~box~~ shows which of the two ~~women~~ he ranks first and ~~which~~ second, and similarly for the ~~girls~~.

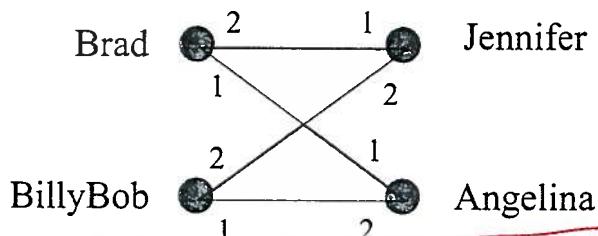


Figure P: Preferences for ~~for~~ four people. ~~Both men like Angelina best and both women like Brad best.~~

More generally, in any matching, ~~a man and a woman who are not married to each other~~ and ~~both men like Angelina best and both women like Brad best.~~

other and who like each other better than their spouses, is called a *rogue couple*. In shown in Figure P,

the situation above, Brad and Angelina would be a rogue couple.

Having a rogue couple is not a good thing, since it threatens the stability of the

marriages. On the other hand, if there are no rogue couples, then for any ~~man and woman~~ who are not married to each other, at least one likes their spouse better than ~~they~~ the other, and so won't be tempted to start an affair.

Definition 9.2.1. A *stable matching* is a matching with no rogue couples.

The question is, given everybody's preferences, how do you find a stable set

In Figure P

of marriages? In the example consisting solely of the four people ~~left~~, we could

let Brad and Angelina both have their first choices by marrying each other. Now

neither Brad nor Angelina prefers anybody else to their spouse, so neither will be

in a rogue couple. This leaves Jen not-so-happily married to Billy Bob, but neither

Jen nor Billy Bob can entice somebody else to marry them, *and so we know there is a stable matching.*

It is something of a surprise that there always is a stable matching among a

men women,

group of ~~boys~~ and ~~girls~~, ~~but there is,~~ and we'll shortly explain why. The surprise

springs in part from considering the apparently similar "buddy" matching prob-

lem. That is, if people can be paired off as buddies, regardless of gender, then

a stable matching *may not* be possible. For example, Figure 9.2 shows a situation

with a love triangle and a fourth person who is everyone's last choice. In this figure

Mergatoid's preferences aren't shown because they don't even matter.

Let's see why there is no stable matching:

A7:

Lemma *There is no stable buddy matching among the four people in Figure 9.2.*

Proof. We'll prove this by contradiction.

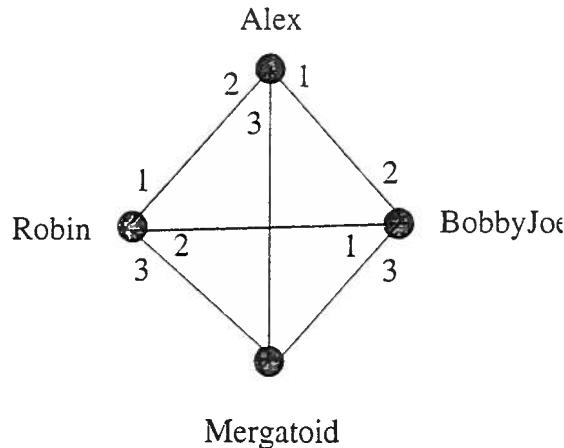


Figure 9.2: Some preferences with no stable buddy matching.

Assume, for the purposes of contradiction, that there is a stable matching. Then there are two members of the love triangle that are matched. Since preferences in the triangle are symmetric, we may assume in particular, that Robin and Alex are matched. Then the other pair must be Bobby-Joe matched with Mergatoid.

But then there is a rogue couple: Alex likes Bobby-Joe best, and Bobby-Joe prefers Alex to his buddy Mergatoid. That is, Alex and Bobby-Joe are a rogue couple, contradicting the assumed stability of the matching. ■

So getting a stable *buddy* matching may not only be hard, it may be impossible.

men

women

But when ~~men~~ are only allowed to marry ~~women~~ and vice versa, then it turns out

*can always be found.*¹
 that a stable matching is not hard to find.

9.2.2 The Mating Ritual

↳ subsection (no #)

The procedure for finding a stable matching involves a *Mating Ritual* that takes

place over several days. The following events happen each day:

woman

Morning: Each ~~girl~~ stands on her balcony. Each ~~boy~~ stands under the balcony

women

of his favorite among the ~~girls~~ on his list, and he serenades her. If a ~~boy~~ has no

women

mother

~~girls~~ left on his list, he stays home and does his ~~homework~~ homework.

woman

Afternoon: Each ~~girl~~ who has one or more suitors serenading her, says to her

favorite among them, "We might get engaged. Come back tomorrow." To the other

suitors, she says, "No. I will never marry you! Take a hike!"

man

woman

woman

Evening: Any ~~girl~~ who is told by a ~~girl~~ to take a hike, crosses that ~~girl~~ off his

list.

when a day arises in which every woman
~~has a suitor~~

Termination condition: When ~~every girl~~ has at most one suitor, the ritual ends

woman

with each ~~girl~~ marrying her suitor, if she has one.

There are a number of facts about this Mating Ritual that we would like to

¹ once again, we disclaim any political statement here — it's just the way that the math works out.

prove:

~~The Ritual eventually reaches the termination condition.~~

- The Ritual has a last day.
- Everybody ends up married.
- The resulting marriages are stable.

9.2.3 A State Machine Model

Before we can prove anything, we should have clear mathematical definitions of what we're talking about. In this section we sketch how to define a rigorous state machine model of the Marriage Problem.

So let's begin by formally defining the problem.

Definition 9.2.2. A *Marriage Problem* consists of two disjoint sets of the same finite size, called the-Boys and the-Girls. The members of the-Boys are called *boys*, and members of the-Girls are called *girls*. For each boy, B , there is a strict total order, \prec_B , on the-Girls, and for each girl, G , there is a strict total order, \prec_G , on the-Boys.

If $G_1 \prec_B G_2$ we say B prefers girl G_2 to girl G_1 . Similarly, if $B_1 \prec_G B_2$ we say G

prefers boy B_2 to boy B_1 .

A *marriage assignment* or *perfect matching* is a bijection, $w : \text{the-Boys} \rightarrow \text{the-Girls}$.

If $B \in \text{the-Boys}$, then $w(B)$ is called B 's *wife* in the assignment, and if $G \in \text{the-Girls}$,

then $w^{-1}(G)$ is called G 's *husband*. A *rogue couple* is a boy, B , and a girl, G , such

that B prefers G to his wife, and G prefers B to her husband. An assignment is

stable if it has no rogue couples. A *solution* to a marriage problem is a stable perfect

matching.

To model the Mating Ritual with a state machine, we make a key observation:

to determine what happens on any day of the Ritual, all we need to know is which

girls are still on which boys' lists on the morning of that day. So we define a state

to be some mathematical data structure providing this information. For example,

we could define a state to be the "still-has-on-his-list" relation, R , between boys

and girls, where $B R G$ means girl G is still on boy B 's list.

We start the Mating Ritual with no girls crossed off. That is, the start state is the

complete bipartite relation in which every boy is related to every girl.

According to the Mating Ritual, on any given morning, a boy will *serenade* the girl he most prefers among those he has not as yet crossed out. Mathematically, the girl he is serenading is just the maximum among the girls on B 's list, ordered by $<_B$. (If the list is empty, he's not serenading anybody.) A girl's *favorite* is just the maximum, under her preference ordering, among the boys serenading her.

Continuing in this way, we could mathematically specify a precise Mating Ritual state machine, but we won't bother. The intended behavior of the Mating Ritual is clear enough that we don't gain much by giving a formal state machine, so we stick to a more memorable description in terms of boys, girls, and their preferences. The point is, though, that it's not hard to define everything using basic mathematical data structures like sets, functions, and relations, if need be.

~~9.2.4 There is a Marriage Day~~

subsubsection (no #)

It's easy to see why the Mating Ritual has a terminal day when people finally get

married. Every day on which the ritual hasn't terminated, at least one ~~boy~~^{man} crosses a ~~girl~~^{woman} off his list. (If the ritual hasn't terminated, there must be some ~~boy~~^{woman} serenaded

by at least two ~~men~~^{men}, and at least one of them will have to cross her off his list).

~~If we start with n men and n women, then~~
~~So starting with n boys and n girls, each of the n boy lists initially has n girls~~

on it, for a total of n^2 list entries. Since no ~~girl~~^{woman} ever gets added to a list, the total

number of entries on the lists decreases every day that the Ritual continues, and so

the Ritual can continue for at most n^2 days.

~~9.2.5~~ They All Live Happily Every After...

↙ Subsubsection (no 2)

We still have to prove that the Mating Ritual leaves everyone in a stable marriage.

To do this, we note one very useful fact about the Ritual: if a ~~girl~~^{woman} has a favorite

~~boy~~^{woman} suitor on some morning of the Ritual, then that favorite suitor will still be

serenading her the next morning —because his list won't have changed. So she is

sure to have today's favorite boy among her suitors tomorrow. That means she will

be able to choose a favorite suitor tomorrow who is at least as desirable to her as

today's favorite. So day by day, her favorite suitor can stay the same or get better,

never worse. In other words, a ~~girl's~~^{woman's} favorite is a weakly increasing variable with

~~respect to her preference order on the boys.~~
 This sounds like an invariant, and it is.

Now we can verify the Mating Ritual using a simple invariant predicate, P .

that captures what's going on.

Let P be the Predicate:

woman, w , man, m , if w has

Definition of P : For every woman w and every boy m , if w is crossed off m 's list, then

w

m

w has a suitor whom she prefers over m .

— INSERT P goes here —

Why is P invariant? Well, we know that w 's favorite tomorrow will be at least

as desirable to her as her favorite today, and since her favorite today is more desir-

able than m , tomorrow's favorite will be too.

Notice that P also holds on the first day, since every w is on every list. So by

words, no one was crossed off any list and

the Invariant Theorem, we know that P holds on every day that the Mating Ritual

P is vacuously true. Hence this means that

runs. Knowing the invariant holds when the Mating Ritual terminates will let us

P is an invariant and thus it must hold on

complete the proof.

on the day when the ritual terminates.

Theorem 9.2.3. Everyone is married by the Mating Ritual.

By Contradiction. Assume

Proof. Suppose, for the sake of contradiction, that it is the last day of the Mating

ritual since there are an equal

Ritual and someone does not get married. Then he can't be serenading anybody,

someone

woman has a suitor

and so his list must be empty. So by invariant P , every woman's favorite boy

since there

— INSERT Q goes here —

INSERT P

P₋)

Lemma 5P : P is an invariant for the Matching Ritual.

Proof : By induction on the number of days.

Base Case : ~~on the first day at the end of day 0,~~
in the beginning (i.e., ~~day 0~~) ~~every~~ woman is on every list — no one has been crossed off and so P is vacuously true.

Inductive Step : Assume P is true ^{at} day d and let w be a woman that has been crossed off ~~on~~ a man m's list ~~by~~ the end of day d+1. If w was crossed off m's list ~~by~~ the end of day d+1, ~~case 1~~: w must have a suitor she prefers on day d+1. Then ~~w was only because~~ ~~w has a suitor~~ prefers on day d+1.

Case 2 : w was crossed off m's list prior to day d+1. ~~By the~~ since P is true ~~at the end of day d~~, this means that w has a suitor she prefers to m on day d. She therefore has the same suitor or someone

she prefers better at the end of day $d+1$. P_{-2}

In both cases, P is true at the end
of day $d+1$ and so P must be an invariant. \blacksquare

~~Using Lemma~~

With Lemma 5 P is hard, we can
now prove:

Thes

INSERT Q

number of ~~men~~ men and women, and since bigamy is not allowed, this means that at least one man (call him Bob) and at least one woman do not get married.

Since Bob is not married, he can't be serenading anybody and so his list must be empty. So by invariant P, every woman has a suitor whom she prefers to Bob. Since it is the last day and every woman still has a suitor, this means that every woman gets married. This ~~is~~ is a contradiction since we already argued that at least one woman is not married. Hence our assumption must be false and so everyone must be married. \square

~~Bob o~~ Since it is the last day, this means that whom she prefers to that boy. In particular, every girl has a favorite boy whom every woman has a suitor that she marries on the last day. So all the girls are married. What's more there is no ruled out, and since there are an equal number of men and women, this means that every man bigny: a boy only serenades one girl, so no two girls have the same favorite. But there are the same number of girls as boys, so all the boys must be married must also be married. But this contradicts the assumption. ■

Theorem 9.2.4. *The Mating Ritual produces a stable matching.*

~~Proof.~~ Let Brad be some boy and Jen be any girl that he is not married to on the last day of the Mating Ritual. We claim that Brad and Jen are not a rogue couple. Since and this that all Brad is an arbitrary boy, it follows that no boy is part of a rogue couple. Hence the marriages on the last day are stable. There are two cases to consider.

To prove the claim, we consider two cases.

by the end.

Case 1. Jen is not on Brad's list. Then by invariant P, we know that Jen prefers (and hence a husband) that she prefers her husband to Brad. So she's not going to run off with Brad. The claim holds in this case. — Brad and Jen has a suitor

cannot be a rogue couple.

Case 2. ~~Otherwise~~, Jen is on Brad's list. But since Brad is not married to Jen, he

must be choosing to serenade his wife instead of Jen, so he must prefer his wife.

- once again, Brad and Jenn

So he's not going to run off with Jen ~~the claim also holds in this case~~. ■

are not a rogue couple -

~~9.2.6~~ ...Especially the ~~men~~  subsection (no #)

~~men~~ ~~women~~? ~~the women~~

Who is favored by the Mating Ritual, the ~~men~~ or the ~~women~~? ~~The girls~~ seem to have

all the power: they stand on their balconies choosing the finest among their suitors

and spurning the rest. What's more, we know their suitors can only change for

~~man~~ ~~woman~~

the better as the Ritual progresses. Similarly, a ~~man~~ keeps serenading the ~~woman~~ he

most prefers among those on his list until he must cross her off, at which point he

~~man's perspective,~~

serenades the next most preferred girl on his list. So from the ~~man's point of view,~~

~~woman~~

the ~~man~~ he is serenading can only change for the worse. Sounds like a good deal

for the ~~man~~ ~~women~~.

~~men~~

But it's not! The fact is that from the beginning, the ~~men~~ are serenading their

~~woman~~

~~woman~~

first choice ~~man~~ and the desirability of the ~~man~~ being serenaded decreases only

~~man~~

~~Mating Ritual~~

enough to give the ~~man~~ his most desirable possible spouse. The ~~mating algorithm~~

~~men~~

actually does as well as possible for all the ~~men~~ and does the worst possible job

women,
for the ~~girls~~

To explain all this we need some definitions. Let's begin by observing that

The Matching Ritual

while the mating algorithm produces one stable matching, there may be other sta-

ble matchings among the same set of ~~boys~~ and ~~girls~~. For example, reversing the

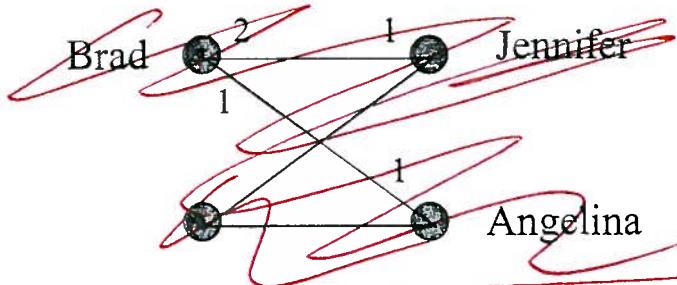
roles of ~~boys~~ and ~~girls~~ will often yield a different stable matching among them.

But some spouses might be out of the question in all possible stable matchings.

Given the preferences shown in Figure P,

For example, Brad is just not in the realm of possibility for Jennifer, since if you

ever pair them, Brad and Angelina will form a rogue couple ~~here's a picture~~



a set of preferences for lists for all men and women

Definition 9.2.5. Given any marriage problem, one person is in another person's

realm of possible spouses if there is a stable matching in which the two people are

married. A person's *optimal spouse* is their most preferred person within their realm

of possibility. A person's *pessimal spouse* is their least preferred person in their

realm of possibility.

Everybody has an optimal and a pessimal spouse, since we know there is at least one stable matching, namely the one produced by the Mating Ritual. Now here is the shocking truth about the Mating Ritual:

men

Theorem 9.2.6. *The Mating Ritual marries every ~~man~~ to his optimal spouse.*

By contradiction.

Proof. Assume for the purpose of contradiction that some ~~man~~ does not get his ~~spouse~~.

Then there must have been a day when he crossed off his optimal ~~girl~~ ~~spouse~~ (and would ultimately marry).

—otherwise he would still be serenading her or some even more desirable girl.

man

By the Well Ordering Principle, there must be a first day when a ~~boy~~ call him

spouse (call her Nicole).

According to the rules of the Ritual, Keith crosses off Nicole because Nicole has

preferred (call him Tom), so

a favorite suitor Tom and

Nicole prefers Tom to Keith (*).

~~Remember this is a proof by contradiction.~~

woman

that

~~Now~~ since this is the first day an optimal ~~man~~ gets crossed off, we know Tom

~~had not previously crossed off his optimal spouse, and so
hasn't crossed off his optimal girl. So~~

Tom ranks Nicole at least as high as his optimal ~~girl~~ (**)

~~spouse~~
~~woman~~

By the definition of an optimal girl, there must be some stable set of marriages in

~~spouse~~

which Keith gets his optimal ~~girl~~, Nicole. But then the preferences given in (*)

and (**) imply that Nicole and Tom are a rogue couple within this supposedly

stable set of marriages (think about it). This is a contradiction. ■

~~women~~

Theorem 9.2.7. *The Mating Ritual marries every ~~girl~~ to her pessimal spouse.*

~~By contradiction. Assume the theorem is not true. QED~~

~~Proof. Say Nicole and Keith marry each other as a result of the Mating Ritual. By
being part of~~

~~the previous Theorem 9.2.6, Nicole is Keith's optimal spouse, and so in any stable~~

~~set of marriages,~~

Keith rates Nicole at least as high as his spouse. (+)

~~Now suppose for the purpose of contradiction that there is another stable set of
marriages where Nicole does worse than Keith. That is, Nicole is married to Tom,~~

and

~~— INSERT R goes here —~~

INSERT R

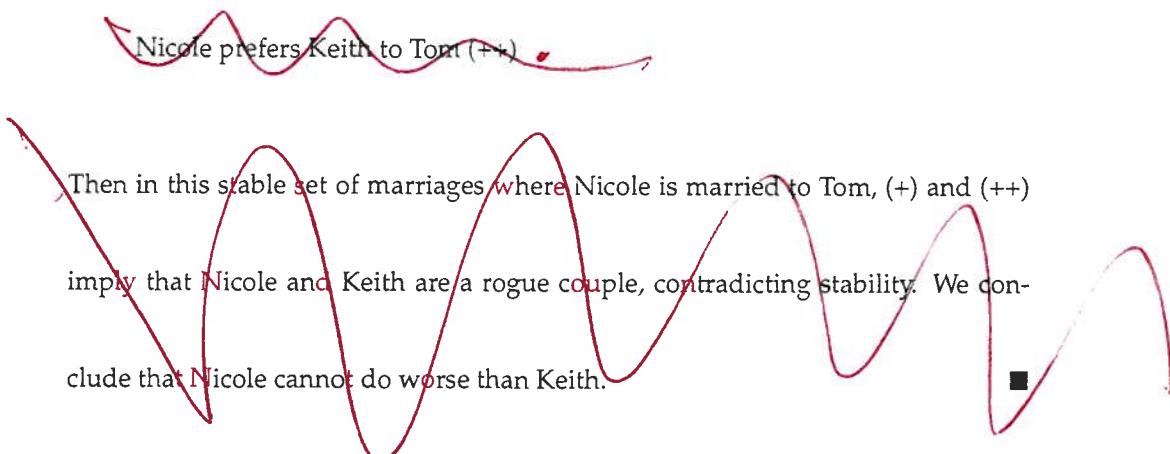
~~and let M' be~~

Hence there must be a stable ^{set of} marriages M' where some woman (call her Nicole) is married to ~~someday~~ a man (call him Tom) that she likes less than her spouse in the Matching Ritual (call him Keith). This means that

Nicole prefers Keith to Tom (+).

By Theorem 9.2.6 and the fact that Nicole and Keith are ~~not~~ married in the Matching Ritual, we know that Keith prefers Nicole to ~~his spouse~~ in M (++)

This means that Keith and Nicole form a rogue couple in M , which contradicts the stability of M . ■



~~9.2.1~~ Applications *↳ subsubsection (note)*

on-line

Not surprisingly, a stable matching procedure is used by at least one large dating

man-woman-

agency. But although "boy-girl-marriage" terminology is traditional and makes

some of the definitions easier to remember (we hope without offending anyone),

solutions to the Stable Marriage Problem are widely useful.

The Mating Ritual was first announced in a paper by D. Gale and L.S. Shapley

in 1962, but ten years before the Gale-Shapley paper was appeared, and unknown

by them, the Ritual was being used to assign residents to hospitals by the National

Resident Matching Program (NRMP). The NRMP has, since the turn of the twen-

tieth century, assigned each year's pool of medical school graduates to hospital

residencies (formerly called "internships") with hospitals and graduates playing

~~men women~~

~~women~~

the roles of ~~boys~~ and ~~girls~~. (In this case there may be multiple ~~boys~~ married to

~~marry a scenario we consider in the Problem Sheet~~

~~one girl, but there's an easy way to use the Ritual in this situation (see Problem ??).~~

~~at the end of the chapter.)~~

Before the Ritual was adopted, there were chronic disruptions and awkward coun-

termeasures taken to preserve assignments of graduates to residencies. The Rit-

ual resolved these problems so successfully, that it was used essentially without

change at least through 1989.²

mating Ritual

The internet infrastructure company, Akamai, also uses a variation of the ~~Gale-Shapley~~

its

~~Shapley procedure~~ to assign web traffic to servers. In the early days, Akamai used

other combinatorial optimization algorithms that got to be too slow as the number

GS,000

of servers (over 20,000 in 2010) ~~became~~ and traffic increased. Akamai ~~increased.~~

and requests (over 800 billion per day)

~~(over 5 terabits per second)~~

switched to ~~Gale-Shapley~~ since it is fast and can be run in a distributed manner.

The mating Ritual

women

correspond to men.

In this case, the ~~web traffic~~ corresponds to the ~~boys~~ and the ~~web servers~~ to the ~~girls~~

we're requests

²Much more about the Stable Marriage Problem can be found in the very readable mathematical

monograph by Dan Gusfield and Robert W. Irving, *The Stable Marriage Problem: Structure and Algorithms*, MIT Press, Cambridge, Massachusetts, 1989, 240 pp.

~~The web requests~~

~~girls. The servers have preferences based on latency and packet loss; the traffic has~~

~~and the web servers have~~

preferences based on the cost of bandwidth, and colocation.

9.2.8 Problems

Practice Problems

Class Problems

Homework Problems

~~TRANSFER SUGGESTIONS
THIS IS THE OLD SECTION. BORROWED FROM IT~~

EDITING NOTE: Add problem proving that the Mating Ritual need not proceed

in morning/afternoon/evening lock step: a girl can reject nonfavorite suitors one

at a time and at any time, and a rejected boy can change the girl he serenades

without waiting for the other boys to change. The proof uses the fact that single

actions commute, so induction proves that all executions are confluent—which

implies all executions end with the same boy-optimal matching. This lemma can

be cited in the planar graphs section to prove that the edges in an embedding can

be added in any order.

More to problem
Geckos at end of
chapter

on pp 414 - 422
—INSERT S (old section 9.5) goes here—

390

CHAPTER 9. SIMPLE GRAPHS

~~9.3 Connectedness~~

5.4 Getting From A to B in a Graph

~~9.4~~

~~5.4.1~~

~~Paths and Simple Cycles~~

~~Paths in simple graphs are essentially the same as paths in digraphs. We just modify the digraph definitions using undirected edges instead of directed ones. For example, the formal definition of a path in a simple graph is virtually that same as Definition 7.1.1 of paths in digraphs:~~

Definition 9.3.1. A *path* in a graph, G , is a sequence of ~~a~~ $\neq 0$ vertices

v_0, \dots, v_k

such that $v_i — v_{i+1}$ is an edge of G for all i where $0 \leq i < k$. The path is said to *start*

at v_0 , to *end* at v_k , and the *length* of the path is defined to be k .

An edge, $u — v$, is *traversed* n times by the path if there are n different values of

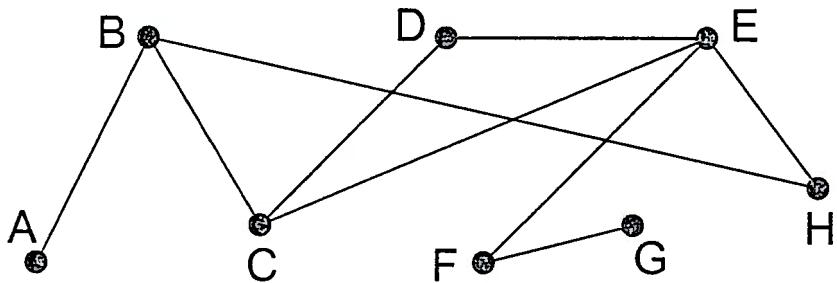
i such that $v_i — v_{i+1} = u — v$. The path is *simple*³ iff all the v_i 's are different, that is,

³Heads up: what we call "paths" are commonly referred to in graph theory texts as "walks," and simple paths are referred to as just "paths". ~~This is what we will call cycles and simple cycles are commonly called "closed walks" and just "cycles."~~

~~$i \neq j$ implies $v_i \neq v_j$.~~

For example, the graph in Figure 9.3 has a length 6 simple path A,B,C,D,E,F,G.

This is the longest simple path in the graph.



*containing a
length 6 simple path*

Figure 9.3: A graph ~~with 5 simple cycles~~

(e.g., A,B,C,D,E,F,G) of length 6.

~~As in digraphs, the length of a path is the total number of times it traverses~~

edges, which is *one less* than its length as a sequence of vertices. For example, the

length 6 path A,B,C,D,E,F,G is actually a sequence of *seven* vertices.

A cycle can be described by a path that begins and ends with the same vertex.

For example, B,C,D,E,C,B is a cycle in the graph in Figure 9.3. This path suggests

that the cycle begins and ends at vertex B, but a cycle isn't intended to have a

*This is
next
insert 5AB
and go to
the new section 5.6*

beginning and end, and can be described by *any* of the paths that go around it. For example, D,E,C,B,C,D describes this same cycle as though it started and ended at D, and D,C,B,C,E,D describes the same cycle as though it started and ended at D but went in the opposite direction. (By convention, a single vertex is a length 0 cycle beginning and ending at the vertex.)

All the paths that describe the same cycle have the same length which is defined to be the *length of the cycle*. (Note that this implies that going around the same cycle twice is considered to be different than going around it once.)

A *simple cycle* is a cycle that doesn't cross or backtrack on itself. For example, the graph in Figure 9.3 has three simple cycles B,H,E,C,B and C,D,E,C and B,C,D,E,H,B. More precisely, a simple cycle is a cycle that can be described by a path of length at least three whose vertices are all different except for the beginning and end vertices. So in contrast to simple *paths*, the length of a simple cycle is the *same* as the number of distinct vertices that appear in it.

From now on we'll stop being picky about distinguishing a cycle from a path

that describes it, and we'll just refer to the path as a cycle.⁴

Simple cycles are especially important, so we will give a proper definition of them. Namely, we'll define a simple cycle in G to be a *subgraph* of G that looks like a cycle that doesn't cross itself. Formally:

Definition 9.3.2. A *subgraph*, G' , of a graph, G , is a graph whose vertices, V' , are a subset of the vertices of G and whose edges are a subset of the edges of G .

Notice that since a subgraph is itself a graph, the endpoints of every edge of G' must be vertices in V' .

Definition 9.3.3. For $n \geq 3$, let C_n be the graph with vertices $1, \dots, n$ and edges

$$1-2, 2-3, \dots, (n-1)-n, n-1.$$

A graph is a *simple cycle* of length n iff it is isomorphic to C_n for some $n \geq 3$. A *simple cycle of a graph*, G , is a subgraph of G that is a simple cycle.

⁴Technically speaking, we haven't ever defined what a cycle is, only how to describe it with paths.

But we won't need an abstract definition of cycle, since all that matters about a cycle is which paths describe it.

This definition formally captures the idea that simple cycles don't have direction or beginnings or ends.

~~9.3.2 Connected Components~~ INSERT 5AC goes here
c.f. is old sec 9.3.4)

9.3.2 Connected Components

S.5 Connectivity

INSERT 5AD goes here
(new text)

Definition 9.3.4. Two vertices in a graph are said to be *connected* when there is

a path that begins at one and ends at the other. By convention, every vertex is

considered to be connected to itself by a path of length zero.

~~EDITING NOTE:~~

Now if there is a path from vertex u to vertex v , then v is connected to u by the reverse path, so connectedness is a symmetric relation. Also, if there is a path from u to v , and also a path from v to w , then these two paths can be combined to form a path from u to w . So the connectedness relation is transitive. It is also reflexive, since every vertex is by definition connected to itself by a path of length zero.

INSERT 5AE moves up to here

INSERT 5AF goes here

~~The~~ diagram in Figure 9.5 looks like a picture of three graphs, but is intended

5.4.3 Numbers of Paths~~Answers
Final answer~~

Given a pair of nodes that are connected by a path ^{of length k in a graph}, there are often many paths ~~that connect~~ of length k on the graph that ~~connect~~ one node to the other. ~~can~~ can be used to get from ~~it~~. For example, there are ~~5~~ paths of length 3 that start at ~~v1~~ and end at ~~v2~~ in the graph shown in Figure AD.

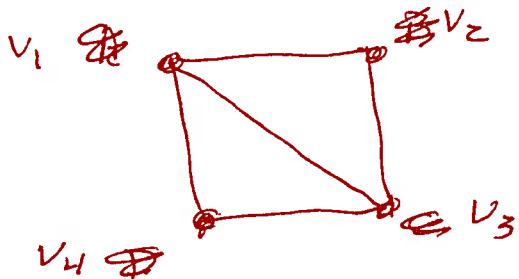


Figure AD: A graph in which there are ~~5~~ paths of length 3 from ~~v1~~ to ~~v2~~. The paths are ~~ABAB, ACAB, ADAB, ADCA~~.
~~(v1, v2, v1, v4), (v1, v3, v1, v4), (v1, v4, v1, v4), (v1, v2, v3, v4), and (v1, v4, v3, v4)~~.
 There is a surprising relationship between the number of paths ~~between two nodes~~ of length k between ~~a pair~~ a pair of nodes in a graph and the k th power of the adjacency matrix A^k for G . The relationship is captured in the following

Theorem.

Theorem 5AE : Let $G = (V, E)$ be an n -node graph with $V = \{v_1, v_2, \dots, v_n\}$, and let $A_G = \{a_{ij}\}$ denote the adjacency matrix for G . Let $a_{ij}^{(k)}$ denote the (i, j) -entry of the k th power of A_G . Then the number of paths of length k between v_i and v_j is $a_{ij}^{(k)}$.

In other words, we can determine the number of paths of length k between any pair of nodes simply by computing the k th power of the adjacency matrix! That's pretty amazing.

For example, the first three powers of the adjacency matrix for the graph in Figure AD are:

$$A^1 = \begin{pmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 0 \\ 1 & 1 & 0 & 1 \\ 1 & 0 & 1 & 0 \end{pmatrix} \quad A^2 = \begin{pmatrix} 3 & 1 & 2 & 2 \\ 1 & 2 & 1 & 2 \\ 2 & 1 & 3 & 1 \\ 1 & 2 & 1 & 2 \end{pmatrix}$$

$$A^3 = \begin{pmatrix} 4 & 5 & 5 & 5 \\ 5 & 2 & 5 & 2 \\ 5 & 5 & 4 & 5 \\ 5 & 2 & 5 & 2 \end{pmatrix}$$

Sure enough, ~~$a_{14}^{(3)}$~~ the $(1, 4)$ coordinate of A^3 is $a_{14}^{(3)} = 5$, which ~~is~~ is the ~~the~~ number of length 3 paths from v_1 to v_4 . And, $a_{24}^{(3)} = 2$, which is the ~~the~~ number of length 3 paths from v_2 to v_4 . ~~Let's~~ By proving the theorem, we'll discover why it is true and thereby uncover the relationship between matrix multiplication and numbers of paths.

Proof of Theorem 5AE: The proof is by induction on k . We will let $P(k)$ be the predicate that the theorem is true ~~for paths of length k~~ for k . ~~We will denote the number of paths between v_i and v_j of length k between v_i and v_j .~~

Base case ($k=1$): ~~There is a path of length~~

~~$k+1$ between v_i and v_j iff there is an edge connecting v_i and v_j . This is true iff $a_{ij}^{(1)} = a_{ij} = 1$. There can be at most one path of length 1 between v_i and v_j . There are 0 paths between v_i and v_j iff $v_i \neq v_j$ and $v_i, v_j \notin E$, which is true iff~~

~~The goal is to show~~

Let $P_{ij}^{(k)}$ denote the number of paths of length k between v_i and v_j . The $P(k)$ is the predicate:

$$\forall i, j \in [1, n]. P_{ij}^{(k)} = \alpha_{ij}^{(k)}. \quad (\text{eqn 5AE})$$

Base Case: ($k=1$). There are two cases to consider.

Case 1: $\{v_i, v_j\} \subseteq E$. Then ~~$P_{ij}^{(1)}$~~ $P_{ij}^{(1)} = 1$ since there is precisely ~~one~~ path of length 1 between v_i and v_j . Moreover, $\{v_i, v_j\} \not\subseteq E$ means that $\alpha_{ij}^{(1)} = \alpha_{ij} = 1$. So $P_{ij}^{(1)} = \alpha_{ij}^{(1)}$ in this case.

Case 2: $\{v_i, v_j\} \not\subseteq E$. Then $P_{ij}^{(1)} = 0$ since there cannot be any paths of length 1 between v_i and v_j . Moreover, $\{v_i, v_j\} \not\subseteq E$ means that $\alpha_{ij}^{(1)} = 0$. So $P_{ij}^{(1)} = \alpha_{ij}^{(1)}$ in this case as well.

Hence, $P(1)$ must be true.

Inductive Step: Assume $P(k)$ is true. In other words, assume that Equation 5AE holds.

We can count the number of paths of length $k+1$ from v_i to v_j according to the ~~first node visited after~~ (call it v_t) ~~after~~
~~the~~ edge in the path (call it $\{v_i, v_t\}$). This means that

$$P_{ij}^{(k+1)} = \sum_{\substack{t: \{v_i, v_t\} \in E \\ t \neq i}} P_{t,j}^{(k)} \quad \text{Eqn AF}$$

where the sum is over all t such that $\{v_i, v_t\}$ is an edge. Using the fact that $a_{it} = 1$ if $\{v_i, v_t\} \in E$ and $a_{it} = 0$ otherwise, we can rewrite Equation AF as follows:

$$P_{ij}^{(k+1)} = \sum_{t=1}^n a_{it} P_{t,j}^{(k)} .$$

By the inductive hypothesis, $P_{t,j}^{(k)} = a_{t,j}^{(k)}$ and thus

$$P_{ij}^{(k+1)} = \sum_{t=1}^n a_{it} a_{t,j}^{(k)} .$$

But the formula for matrix multiplication gives that

$$a_{i,j}^{(k+1)} = \sum_{t=1}^n a_{i,t} a_{t,j}^{(k)}$$

and so we must have $p_{i,j}^{(k+1)} = a_{i,j}^{(k+1)}$

for all $i, j \in [1, n]$. Hence $P(k+1)$ is true and the induction is complete. \blacksquare

5.4.4 Shortest Paths

Although the connection between the power of the adjacency matrix and the number of paths is cool (at least ~~if you're a mathematician~~, ~~cool~~ if you are a mathematician), ~~counting~~ the problem of counting paths does not come up very often in practice. Much more important is the problem of finding the shortest path ~~in a graph~~ between a ~~pair~~ pair of nodes in a graph. There is good news and bad news to report on this front.

The good news is that it is not very hard to find a shortest path. ~~there are several ways to do it~~ The bad news is that you can't win one of those million dollar prizes for doing it.

In fact, there are several good algorithms known for finding ~~the~~ a shortest path between a pair of nodes. The simplest to explain ~~in fact, one way~~ (but not ~~the best~~) is to compute the powers of the adjacency matrix one by one until the value of $a_{ij}^{(k)}$ exceeds 0. That's because $a_{ij}^{(k)}$ ~~is the smallest value by theorem~~ implies that ~~the length of the shortest path between v_i and v_j will be the smallest value of k for which $a_{ij}^{(k)} > 0$.~~

~~Two~~ Paths on Weighted Graphs \Leftarrow ^{subsubsection} _(not #)

~~The scene~~

The problem of computing shortest paths in a weighted graph frequently arises in practice. For example, when you drive home

for vacation, you usually would like to take the shortest route. ~~In the case of a weighted graph, the length of a path is simply the sum of the weights~~

Definition 5H: Given a weighted graph, the length of a path in the graph is the sum of the weights of the edges in the path.

Finding shortest paths in weighted graphs is not a lot harder than finding shortest paths in unweighted graphs. We won't show you how to do it here, but ~~Induction once again plays~~ ^{finding} you will ~~see~~ ^{if} study algorithms for shortest paths ~~when~~ you take an algorithms course. Not surprisingly, the proof of correctness will use induction.

INSERT 5 AF

~~DATA~~

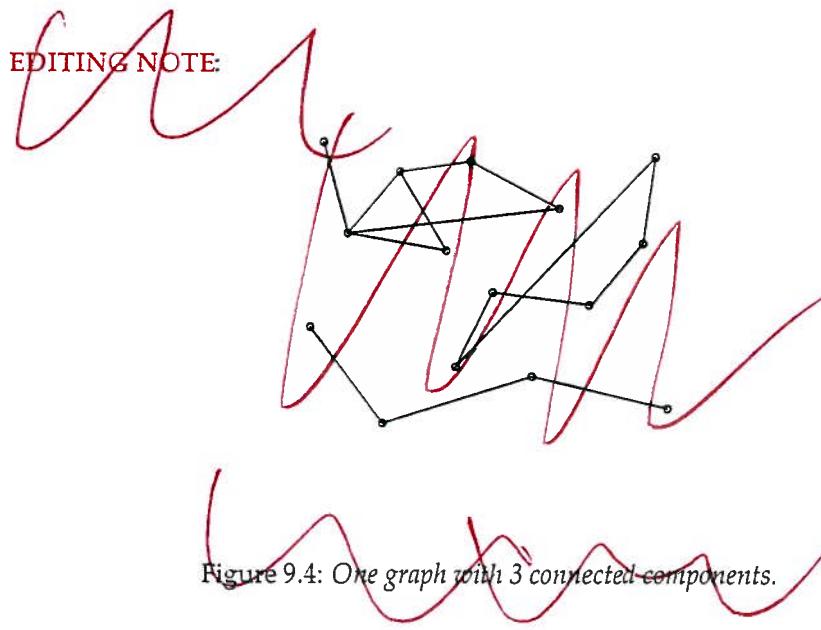
5.5.1 Connected Components

~~GOALS~~

Being connected is usually a good property for a graph to have. For example, it could mean that it is possible to get from any node to any other node, ~~box~~ or that it is possible to communicate between any pair of nodes, depending on the application.

But not all graphs are connected. For example, the graph ~~represents~~ where nodes represent cities and ~~whose~~ edges represent highways might be connected for North American cities, but would ~~not~~ surely not be connected if you ^{else} included cities in Australia. The same is true for communication networks like the Internet — in order to be ~~protected~~ from viruses that spread on the Internet, some government networks are completely isolated from the Internet. For example, the

to be a picture of *one* graph. This graph consists of three pieces (subgraphs). Each piece by itself is connected, but there are no paths between vertices in different pieces.



To move
beginning
of S.S

This is my
S.S E

Definition 9.3.5. A graph is said to be *connected* when every pair of vertices are connected.

~~These connected pieces of a graph are called its *connected components*.~~

DEFINITION 54 : A

~~Our definition is easy: a connected component is the set of all the vertices connected~~

~~is a subgraph of a~~
~~graph consisting of some vertex and every node~~
~~and edge that is connected to that vertex.~~



Figure 9.5: One graph with 3 connected components.

~~connected to some single vertex~~ So a graph is connected iff it has exactly one connected component.

The empty graph on n vertices has n connected components.

5.5.2 k -connected graphs

~~9.3.3 How Well Connected?~~

If we think of a graph as modelling cables in a telephone network, or oil pipelines,

or electrical power lines, then we not only want connectivity, but we want connectivity that survives component failure. A graph is called *k-edge connected* if it takes

at least k "edge-failures" to disconnect it. More precisely:

Definition 9.3.6. Two vertices in a graph are *k-edge connected* if they remain con-

nected in every subgraph obtained by deleting $k - 1$ edges. A graph with at least

two vertices is k -edge connected⁵ if every two of its vertices are k -edge connected.

So 1-edge connected is the same as connected for both vertices and graphs. Another way to say that a graph is k -edge connected is that every subgraph obtained from it by deleting at most $k - 1$ edges is connected. For example, in the graph in *vertices C and E are 3-edge connected,*
 Figure 9.3 ~~A~~⁴ vertices B and E are 2-edge connected, G and E are 1-edge connected,
 and no vertices are ~~2~~-edge connected. The graph as a whole is only 1-edge connected. More generally, any simple cycle is 2-edge connected, and the complete graph, K_n , is $(n - 1)$ -edge connected.

If two vertices are connected by k edge-disjoint paths (that is, no two paths traverse the same edge), then they are obviously k -edge connected. A fundamental fact, whose ingenious proof we omit, is Menger's theorem which confirms that the converse is also true: if two vertices are k -edge connected, then there are k edge-disjoint paths connecting them. It even takes some ingenuity to prove this for the

⁵The corresponding definition of connectedness based on deleting vertices rather than edges is common in Graph Theory texts and is usually simply called " k -connected" rather than " k -vertex connected."

case $k = 2$.

~~5.4.2 Finding The 9.3.4 Connection by Simple Path~~

~~E32~~

Where there's a path, there's a simple path. This is sort of obvious, but it's easy

enough to prove rigorously using the Well Ordering Principle.

Lemma 9.3.7. *If vertex u is connected to vertex v in a graph, then there is a simple path from u to v .*

Proof. Since there is a path from u to v , there must, by the Well-ordering Principle,

be a minimum length path from u to v . If the minimum length is zero or one, this

minimum length path is itself a simple path from u to v . Otherwise, there is a

minimum length path

$$v_0, v_1, \dots, v_k$$

from $u = v_0$ to $v = v_k$ where $k \geq 2$. We claim this path must be simple. To

prove the claim, suppose to the contrary that the path is not simple, that is, some

vertex on the path occurs twice. This means that there are integers i, j such that

THIS IS INSERT SAC and more after sections. 4.8!

$0 \leq i < j \leq k$ with $v_i = v_j$. Then deleting the subsequence

$$v_{i+1}, \dots, v_j$$

yields a strictly shorter path

$$v_0, v_1, \dots, v_i, v_{j+1}, v_{j+2}, \dots, v_k$$

from u to v , contradicting the minimality of the given path. ■

Actually, we proved something stronger:

Corollary 9.3.8. *For any path of length k in a graph, there is a simple path of length at most k with the same endpoints.*

5.5.3

~~9.3.5~~ The Minimum Number of Edges in a Connected Graph

~~In order to be~~

The following theorem says that a graph with few edges must have many connected components.

Theorem 9.3.9. *Every graph with v vertices and e edges has at least $v - e$ connected components.*

Of course for Theorem 9.3.9 to be of any use, there must be fewer edges than vertices.

Proof. We use induction on the number of edges, e . Let $P(e)$ be the proposition that

for every v , every graph with v vertices and e edges has at least $v - e$ connected components.

Base case: ($e = 0$). In a graph with 0 edges and v vertices, each vertex is itself a connected component, and so there are exactly $v = v - 0$ connected components.

So $P(e)$ holds.

Inductive step: Now we assume that the induction hypothesis holds for every e -edge graph in order to prove that it holds for every $(e + 1)$ -edge graph, where

$e \geq 0$. Consider a graph, G , with $e + 1$ edges and \check{v} vertices. We want to prove that

G has at least $v - (e + 1)$ connected components. To do this, remove an arbitrary edge $a - b$ and call the resulting graph G' . By the induction assumption, G' has at least $v - e$ connected components. Now add back the edge $a - b$ to obtain the

original graph G . If a and b were in the same connected component of G' , then G

has the same connected components as G' , so G has at least $v - e > v - (e + 1)$

components. Otherwise, if a and b were in different connected components of G' ,

then these two components are merged into one in G , but all other components

remain unchanged, reducing the number of components by 1. Therefore, G has at

least $(v - e) - 1 = v - (e + 1)$ connected components. So in either case, $P(e+1)$ holds.

This completes the Inductive step. The theorem now follows by induction. ■

Corollary 9.3.10. *Every connected graph with v vertices has at least $v - 1$ edges.*

A couple of points about the proof of Theorem 9.3.9 are worth noticing. First,

we used induction on the number of edges in the graph. This is very common

in proofs involving graphs, and so is induction on the number of vertices. When

you're presented with a graph problem, these two approaches should be among

the first you consider. The second point is more subtle. Notice that in the inductive

step, we took an arbitrary $(n + 1)$ -edge graph, threw out an edge so that we could

apply the induction assumption, and then put the edge back. You'll see this shrink-

down, grow-back process very often in the inductive steps of proofs related to

graphs. This might seem like needless effort; why not start with an n -edge graph

and add one more to get an $(n + 1)$ -edge graph? That would work fine in this

case, but opens the door to a nasty logical error called *buildup error*, illustrated in

~~Problems ?? and ?? Always use shrink-down, grow-back arguments, and you'll~~

~~never fall into this trap.~~

~~9.3.6 Problems~~

~~Class Problems~~

~~Homework Problems~~

~~Homework Problems~~

~~9.4 Trees~~

- ~~INSERT 5A6 goes here —~~
- ~~INSERT 5A4 goes here —~~
- ~~INSERT 5A1 goes here —~~
(text from end of old
Section 9.4 on
pp 411-)
- ~~INSERT 5AJ goes here —~~
- ~~INSERT 5AK goes here —~~
(text from ~~pp 413~~)
- ~~INSERT 5AL goes here —~~

Trees are a fundamental data structure in computer science, and there are many

kinds, such as rooted, ordered, and binary trees. In this section we focus on the

INSERT 5A6Graph Theory II5.5.4 1 Build-Up Error

~~Here is a false proof about connectivity. It exposes a very common flaw made on proofs by induction on graphs – it even has a name – it is known as “build-up error”.~~

False Claim. *If every vertex in a graph has degree at least 1, then the graph is connected.*

There are many counterexamples; ~~here is one~~ *for example, see Figure 5Z.*

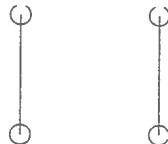


Figure 5Z: A counterexample graph to the False Claim.

Since the claim is false, there must be at least one error in the following “proof”.

False Proof. We use induction. Let $P(n)$ be the proposition that if every vertex in an n -vertex graph has degree at least 1, then the graph is connected.

Base case: There is only one graph with a single vertex and it has degree 0. Therefore, $P(1)$ is vacuously true, since the if-part is false.

Inductive step: We must show that $P(n)$ implies $P(n+1)$ for all $n \geq 1$. Consider an n -vertex graph in which every vertex has degree at least 1. By the assumption $P(n)$, this graph is connected; that is, there is a path between every pair of vertices. Now we add one more vertex x to obtain an $(n+1)$ -vertex graph ~~as shown in Figure 5Y.~~

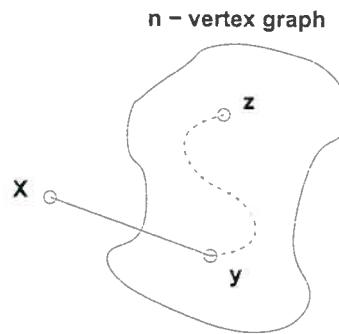


Figure 5Y: Adding a vertex x with degree at least 1 to an n -node graph.

All that remains is to check that there is a path from x to every other vertex z . Since x has degree at least one, there is an edge from x to some other vertex; call it y . Thus, we can obtain a path from x to z by adjoining the edge $x-y$ to the path from y to z . This proves $P(n+1)$.

By the principle of induction $P(n)$ is true for all $n \geq 1$, which proves the theorem. \square

~~The~~ That looks fine! Where is the bug? It turns out that faulty assumption underlying this argument is that *every $(n+1)$ -vertex graph with minimum degree 1 can be obtained from an n -vertex graph with minimum degree 1 by adding 1 more vertex*. Instead of starting by considering an arbitrary $(n+1)$ -node graph, this proof only considered ~~a~~ $(n+1)$ -node graph that you can make by starting with an n -node graph with minimum degree 1.

~~The~~ The counterexample above shows that this assumption is false; there is no way to build ~~that~~ 4-vertex graph from a 3-vertex graph with minimum degree 1. Thus the first error in the proof is the statement "This proves $P(n+1)$ ".

~~This kind of flaw is known as usually build-up error~~ More generally, this is an example of "build-up error". Generally, this arises from a faulty assumption that every size $n+1$ graph with some property can be "built up" from a size n graph with the same property. (This assumption is correct for some properties, but incorrect for others—such as the one in the argument above.)

One way to avoid an accidental build-up error is to use a "shrink down, grow back" process in the inductive step: start with a size $n+1$ graph, remove a vertex (or edge), apply the inductive hypothesis $P(n)$ to the smaller graph, and then add back the vertex (or edge) and argue that $P(n+1)$ holds. Let's see what would have happened if we'd tried to prove the claim above by this method:

Revised

Inductive step: We must show that $P(n)$ implies $P(n+1)$ for all $n \geq 1$. Consider an $(n+1)$ -vertex graph G in which every vertex has degree at least 1. Remove an arbitrary vertex v , leaving an n -vertex graph G' in which every vertex has degree... uh-oh!

The reduced graph G' might contain a vertex of degree 0, making the inductive hypothesis $P(n)$ inapplicable! We are stuck—and properly so, since the claim is false!

~~Always use shrink-down, grow-back arguments and you'll never fall into this trap.~~

2 Trees

Today we'll talk about a very special class of graphs called *trees*. Trees arise in all sorts of applications and you'll see them in just about every computer science class that you'll take at MIT. There are at least half a dozen ways to define a tree, but the simplest is the following.

Definition. A tree is a simple¹, connected, acyclic graph.

~~Recall that we only consider simple graphs in this class; that is, graphs without loops or multiple edges.~~

5.6 Around and Around We Go!

5.6.1 Cycles

Definition¹: ~~The~~ Given a graph $G = (V, E)$, a cycle in G is a sequence of ~~not~~ nodes v_0, v_1, \dots, v_k where ~~such that~~ $v_i - v_{i+1} \in E$ for $0 \leq i < k$ and $v_0 = v_k$. The length of the cycle is k . A cycle is a simple cycle if $k \geq 3$ and ~~the~~ v_0, v_1, \dots, v_{k-1} are all different.

For example, B, C, D, E, C, B is a cycle, ~~on the graph shown in Figure 3.~~ ~~But,~~ it is not simple since it contains node C twice.
There are many ways to represent a cycle. For

¹ Some texts ~~will~~ use the term closed walk or circuit instead of cycle and cycle instead of simple cycle.

Example, B, C, D, E, C, B is the same as C, D, E, C, B, C (just starting at node C instead of node B) ~~isn't~~ and the same as B, C, E, D, C, B (just reversing ~~the~~ direction).

Cycles are very similar to paths, except that the last node is the first node and the notion of first and last does not matter.

5.6.2 Euler Tours

INSERT 5AJ

IT IS NOT ~~difficult~~ difficult to extend Theorem 9.4.3 to prove that a connected graph G has an Euler path if and only if ~~either~~ ^{precisely} 0 or 2 nodes in G have odd degree. Hence, we can conclude that the graph shown in Figure 5BC has an Euler path but not an Euler ~~tour~~ since the graph has precisely two nodes with odd degree.

Although the proof of Theorem 9.4.3 does not explicitly deline a method for finding an Euler tour when one exists, it is not hard to modify the proof to ~~read~~ produce such ~~one~~ a method. The idea is to ~~keep~~ grow ~~up~~ a tour by continually splicing in ~~other~~ cycles until all the edges are consumed.

~~5.6.3 The Travelling Salesman Problem~~
5.6.3 Hamiltonian Cycles and Tours

~~Travelling Salesman Problem~~

~~Travelling Salesperson Problem~~

[#] Hamiltonian cycles are the unruly cousins of the Euler trees.

~~Definition~~ : A Hamiltonian cycle in a graph G is a cycle that visits every node in G exactly once. ~~A~~ ^{similarly,} Hamiltonian path is a path ~~in~~ in G that visits every node exactly once.

Although Hamiltonian cycles sound similar to Euler trees — one visits every node once while the other visits every edge once — ~~they're~~ finding a Hamiltonian cycle can be a lot harder than finding an Euler tree. The same is true for Hamiltonian paths. This is because ~~there is no one~~ ~~known~~ has discovered a

5.6.4 The Travelling Salesperson Problem5.6.4 The Travelling Salesperson Problem

As if the problem of finding a Hamiltonian cycle is not hard enough, when the ~~graph~~ graph is weighted, we often want to find a Hamiltonian cycle that has least possible weight. This is a very famous optimization ~~prob~~ problem known as the Travelling Salesman Problem.

Definition:

~~Definition~~: Given a weighted graph G , the weight of a cycle in G is defined to be the sum of the weights of the edges in the cycle. ~~The Travelling Salesman Problem~~ ~~The problem of finding a Hamiltonian cycle~~

For example, consider the graph shown in Figure AL and suppose that you would like to visit every node once and finish at the node where you started. Can you find a way to do this by traversing a cycle with weight 15?

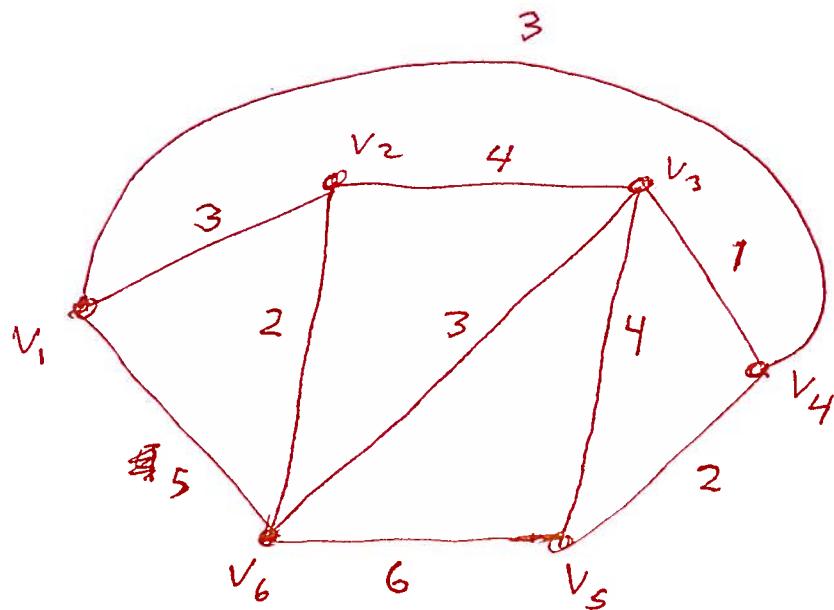


Figure AL : A weighted graph. Can you find a cycle that visits every node exactly once with weight 15?

~~Perhaps it is~~

It needless to say, if you can programme out a fast procedure ~~to always~~ that finds the ~~minimum weight~~ the optimal ~~cycle~~ cycle for the ~~the~~ travelling sales man, you can win a million dollars.[†]

~~I am not the~~

5.7 Trees

~~Cycles~~ As we have just seen, ~~Cycles~~ ^{fundamental}
~~Trees~~ good cycles in a graph can be tricky
 than you might first think. But what if
 a graph has no cycles at all? Sounds
 pretty dull. But ~~in fact~~ graphs without
 cycles (called acyclic graphs) are probably
 the most important graphs of all when
 it comes to computer science. &

~~⇒ Connected~~

Definition SBL: A ~~connected~~ connected
 acyclic graph is called a tree.

Trees are a fundamental data
 structure in computer science and you
 will encounter them in every CS course
 that you take.

purest kind of tree. Namely, we use the term *tree* to mean a connected graph without simple cycles.

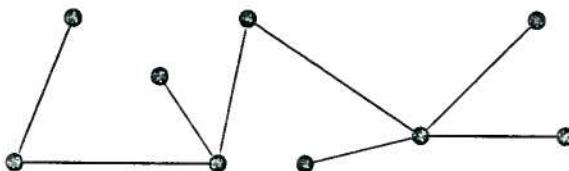
A graph with no simple cycles is called *acyclic*; so trees are acyclic connected graphs.

~~5.7.1~~

~~9.4.1 Tree Properties~~

(STILL TO BE EDITED)

Here is an example of a tree:



A vertex of degree at most one is called a *leaf*. In this example, there are 5 leaves.

Note that the only case where a tree can have a vertex of degree zero is a graph with a single vertex.

The graph shown above would no longer be a tree if any edge were removed, because it would no longer be connected. The graph would also not remain a tree

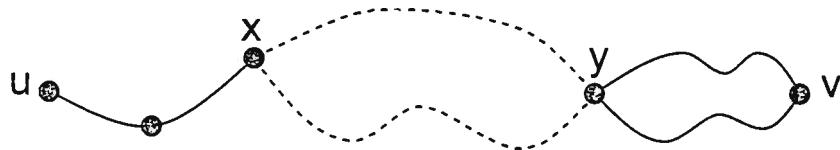
if any edge were added between two of its vertices, because then it would contain a simple cycle. Furthermore, note that there is a unique path between every pair of vertices. These features of the example tree are actually common to all trees.

Theorem 9.4.1. *Every tree has the following properties:*

1. *Any connected subgraph is a tree.*
2. *There is a unique simple path between every pair of vertices.*
3. *Adding an edge between two vertices creates a cycle.*
4. *Removing any edge disconnects the graph.*
5. *If it has at least two vertices, then it has at least two leaves.*
6. *The number of vertices is one larger than the number of edges.*

Proof. 1. A simple cycle in a subgraph is also a simple cycle in the whole graph, so any subgraph of an acyclic graph must also be acyclic. If the subgraph is also connected, then by definition, it is a tree.

2. There is at least one path, and hence one simple path, between every pair of vertices, because the graph is connected. Suppose that there are two different simple paths between vertices u and v . Beginning at u , let x be the first vertex where the paths diverge, and let y be the next vertex they share. Then there are two simple paths from x to y with no common edges, which defines a simple cycle. This is a contradiction, since trees are acyclic. Therefore, there is exactly one simple path between every pair of vertices.



3. An additional edge $u—v$ together with the unique path between u and v forms a simple cycle.
4. Suppose that we remove edge $u—v$. Since the tree contained a unique path between u and v , that path must have been $u—v$. Therefore, when that edge is removed, no path remains, and so the graph is not connected.

5. Let v_1, \dots, v_m be the sequence of vertices on a longest simple path in the tree. Then $m \geq 2$, since a tree with two vertices must contain at least one edge. There cannot be an edge $v_1—v_i$ for $2 < i \leq m$; otherwise, vertices v_1, \dots, v_i would form a simple cycle. Furthermore, there cannot be an edge $u—v_1$ where u is not on the path; otherwise, we could make the path longer. Therefore, the only edge incident to v_1 is $v_1—v_2$, which means that v_1 is a leaf. By a symmetric argument, v_m is a second leaf.

6. We use induction on the number of vertices. For a tree with a single vertex, the claim holds since it has no edges and $0 + 1 = 1$ vertex. Now suppose that the claim holds for all n -vertex trees and consider an $(n+1)$ -vertex tree, T . Let v be a leaf of the tree. You can verify that deleting a vertex of degree 1 (and its incident edge) from any connected graph leaves a connected subgraph. So by 1., deleting v and its incident edge gives a smaller tree, and this smaller tree has one more vertex than edge by induction. If we re-attach the vertex, v , and its incident edge, then the equation still holds because the number of

vertices and number of edges both increase by 1. Thus, the claim holds for T and, by induction, for all trees.

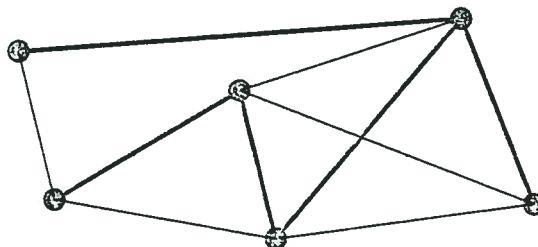
■

Various subsets of these properties provide alternative characterizations of trees, though we won't prove this. For example, a *connected* graph with a number of vertices one larger than the number of edges is necessarily a tree. Also, a graph with unique paths between every pair of vertices is necessarily a tree.

9.7.2

9.4.2 Spanning Trees (SYLL TO B&E EDITED)

Trees are everywhere. In fact, every connected graph contains a subgraph that is a tree with the same vertices as the graph. This is called a *spanning tree* for the graph. For example, here is a connected graph with a spanning tree highlighted.



Theorem 9.4.2. *Every connected graph contains a spanning tree.*

Proof. Let T be a connected subgraph of G , with the same vertices as G , and with

the smallest number of edges possible for such a subgraph. We show that T is acyclic by contradiction. So suppose that T has a cycle with the following edges:

$$v_0—v_1, v_1—v_2, \dots, v_n—v_0$$

Suppose that we remove the last edge, $v_n—v_0$. If a pair of vertices x and y was joined by a path not containing $v_n—v_0$, then they remain joined by that path. On the other hand, if x and y were joined by a path containing $v_n—v_0$, then they remain joined by a path containing the remainder of the cycle. So all the vertices of G are still connected after we remove an edge from T . This is a contradiction, since T was defined to be a minimum size connected subgraph with all the vertices of G . So T must be acyclic. ■

EDITING NOTE:

Tree Variations

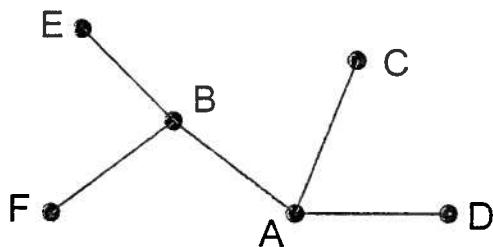
Trees come up often in computer science. For example, information is often stored

in tree-like data structures and the execution of many recursive programs can be

regarded as a traversal of a tree. There are many varieties of trees. For example, a

rooted tree is a tree with one vertex identified as the *root*. Let $u—v$ be an edge in a

rooted tree such that u is closer to the root than v . Then u is the *parent* of v , and v is a

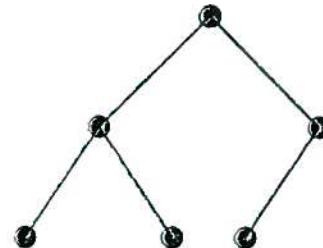


child of u .

In the tree above, suppose that we regard vertex A as the root. Then E and F are

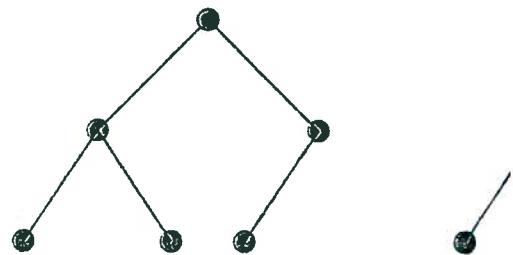
the children of B , and A is the parent of B , C , and D . A *binary tree* is a rooted tree

in which every vertex has at most two children. Here is an example, where the top-



most vertex is the root.

In an *ordered, binary tree*, the children of a vertex v are distinguished. One is called the *left child* of v , and the other is called the *right child*. For example, if we regard the two binary trees below as unordered, then they are equivalent. However, if we re-



gard these trees as ordered, then they are different.

b) Lc

5.7.3 Minimum Weight Spanning Trees (TBA)

EDITING NOTE:

Problem ?? presents most of this as a homework problem.

5.8 Planar Graphs

as 5.8
CHAPTER 11 will go here when I edit it

5.9 Problems

and goes top 402

Devol:
take out
of staff
notes
mode

TRAVERSING A GRAPH

411

~~Traversing a Graph~~

THIS IS INSERT 5AI

Can you walk every hallway in the Museum of Fine Arts *exactly once*? If we represent hallways and intersections with edges and vertices, then this reduces to a question about graphs. For example, could you visit every hallway exactly once in

a museum with ~~this~~ floorplan? *The* ~~this~~ is shown on Figure 5BC?

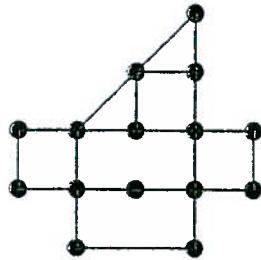


Figure 5BC : A possible floorplan for a museum. Can you find a cycle that traverses every edge exactly once?

The entire field of graph theory began when Euler asked whether the seven bridges

of Königsberg could all be traversed exactly once—essentially the same question

we asked about the Museum of Fine Arts. In his honor, an *Euler path* is a defined

to be a path that traverses every edge in a graph exactly once. Similarly, an *Euler*

tour is an Euler walk that starts and finishes at the same vertex, that is a cycle that

traverses every edge exactly once. Graphs with Euler tours and Euler ~~walks~~^{paths} both have simple characterizations.

connected if and only if

Theorem 9.4.3. A graph has an Euler tour ~~if it is connected and~~ every vertex has even degree.

~~→ INSERT 5CC goes here ←~~

~~Proof. Suppose a graph has an Euler tour. Every pair of vertices must appear in the tour, so the graph is connected. Moreover, a vertex that appears k times in the tour must have degree $2k$, so every vertex of the graph has even degree.~~

~~Unconvincing!~~

~~Conversely, suppose every vertex in a graph G has even degree. Let V_0~~

~~1~~

~~no edge more than once.~~

~~be the longest path in G that traverses every edge at most once. Now~~

~~W must traverse every edge incident to v_0 ; otherwise, the path could be extended and W would not be the longest path that traverses every edge at most once. Moreover, it must be that $v_k = v_0$ and in particular, the W traverses two of these edges each time it passes through v_m . That W is a cycle, since otherwise v_k would have odd degree in W (and hence in G) which is not possible by assumption, and it traverses $v_{m+1} - v_m$ at the end. This accounts for an odd number of edges,~~

~~but the degree of v_m is even by assumption. Therefore, the W must also begin at v_m ; that is, $v_0 = v_m$.~~

~~we conclude the argument with a proof by contradiction.~~

~~Suppose that W is not an Euler tour. Because G is a connected~~

~~1 And you notice that we are using the well~~

~~ordering principle or ~~discrete~~ here? when~~

~~we implicitly assume that a longest ~~path~~ exists?~~

~~This is OK since the length of a ~~path~~ where no edge~~

~~a variation of the~~

~~is used more than once is at most $|E|$.~~

INSERT 5CC

Proof : We first show that if a graph has an Euler tour, then every vertex has even degree.

Assume that a graph $G = (V, E)$ has an Euler tour v_0, v_1, \dots, v_k where $v_k = v_0$.

Since every edge is traversed once in the tour, ~~the degree~~ $k = |E|$ and the degree of a node v_i in G is ~~is twice the number of times that~~ node appears in the sequence $v_0, v_1, \dots, v_{k-1} \in$ times two. We multiply by two since ~~every time~~ ~~is a~~

~~the node appears in the sequence, there is an edge~~ ~~where v_i = v_k~~
in $v = v_i$ for some i , then both ~~v_{i-1}-v_i~~ and $v_i - v_{i+1}$ are edges ~~are~~ incident to v in G . ~~If i=0, then v₀ is replaced by v_k in G.~~ If $v = v_0 = v_k$, then both $v_{k-1}-v_k$ and v_0-v_1 are edges incident to v in G . ~~Note that no edge is ever counted twice in this argument.~~

Hence, the ~~degree of~~ every node is even.

~~Since a~~ we next show that if ~~the degree of every edge node~~ node is even ~~then~~ in a graph $G = (V, E)$, then there is an Euler tour. Let

$$W ::= v_0, v_1, \dots, v_k$$

graph, we can find an edge not in W but incident to some vertex in W . Call this

edge $u-v_i$. But then we can construct a ~~longer walk~~ but that still uses no edge more than once;

$W' := \underline{v} = u, v_i, v_{i+1}, \dots, v_k, \leftarrow v_1, v_2, \dots, v_i \rightarrow$

~~$v, u - v_{i+1}, v_i - v_{i+1}, \dots, v_{k-1} - v_k, u - v_0 - v_1, \dots, v_{k-1} - v_k - v_0 - v_1 - \dots - v_{k-1}$~~

This contradicts the definition of W , so W must be an Euler tour after all. ■

Corollary 9.4.4. A connected graph has an Euler walk if and only if either 0 or 2 vertices

have odd degree.

~~✓ ok as was~~

~~cycles~~ Hamiltonian cycles are the unruly cousins of Euler tours. A ~~Hamiltonian cycle~~

~~is walk that starts and ends at the same vertex and visits every vertex in a graph~~

~~exactly once. That is, no simple characterization of all graphs with a Hamiltonian~~

~~cycle. In fact, determining whether a given graph has a Hamiltonian cycle is~~ ~~H~~

the same category of problem as the SAT problem of Section 1.4: you get a million

dollars for finding an efficient way to determine when a graph has a Hamiltonian

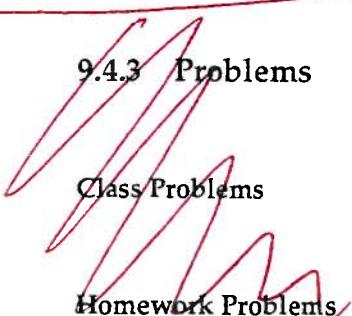
cycle—or for proving that no procedure works efficiently on all graphs.

~~AN~~
This is INSERT 5AK and goes to page 402

This is Insert S and it becomes Section 5.3

414

CHAPTER 9. SIMPLE GRAPHS



5.3 ~~9.5~~ Coloring Graphs

9.2,

In section ~~9.1~~, we used edges to indicate an affinity between two nodes but have considered situations where ~~we need~~ it is useful to use ~~edges to~~ ~~an edge~~ represent a conflict between two nodes also turns out to be really useful.

a pair of nodes. We now consider situations where ~~we need~~ it is useful to use ~~edges to~~ ~~an edge~~ represent a conflict between two nodes also turns out to be really useful.

a pair of nodes. For example, consider the following exam scheduling problem.

9.6 Modelling Scheduling Conflicts

5.3.1 ~~S.T.~~ An Exam Scheduling Problem

Each term, the MIT Schedules Office must assign a time slot for each final exam.

This is not easy, because some students are taking several classes with finals, and (even at MIT)

a student can take only one test during a particular time slot. The Schedules Office

wants to avoid all conflicts. Of course, you can make such a schedule by having

every exam in a different slot, but then you would need hundreds of slots for the

hundreds of courses, and exam period would run all year! So, the Schedules Office

would also like to keep exam period short. The Schedules Office's problem is easy

to describe as a graph. There will be a vertex for each course with a final exam, and

two vertices will be adjacent exactly when some student is taking both courses.

For example, suppose we need to schedule exams for 6.041, 6.042, 6.002, 6.003 and

6.170. The scheduling graph might look like this:

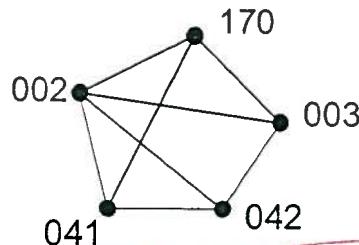


Figure 5R : A scheduling graph for five exams.
Exams connected by an edge cannot be given at the same time.

both courses, so there is an edge between their nodes. On the other hand, 6.042 and

6.170 can have an exam at the same time if they're taught at the same time (which

they sometimes are), since no student can be enrolled in both (that is, no student

should be enrolled in both when they have a timing conflict). Next, identify each
we next identify each

time slot with a color. For example, Monday morning is red, Monday afternoon is

blue, Tuesday morning is green, etc.

Then
 Assigning an exam to a time slot is ~~not~~ equivalent to coloring the corresponding vertex.

The main constraint is that *adjacent vertices must get different colors* —

otherwise, some student has two exams at the same time. Furthermore, in order

to keep the exam period short, we should try to color all the vertices using as *few*

As shown in Figure 5S, different colors as possible. For our example graph, three colors suffice: for our example.

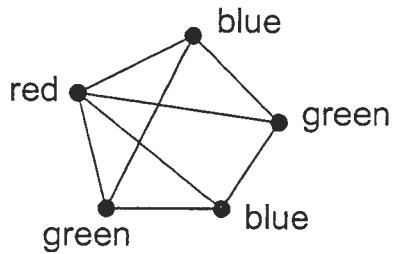


Figure 5S! A 3-coloring of the exam graph from Figure 5R.

This coloring corresponds to giving one final on Monday morning (red), two

~~in Figure 5~~

Monday afternoon (blue), and two Tuesday morning (green). Can we use fewer

than three colors? No! We can't use only two colors since there is a triangle in the

graph, and three vertices in a triangle must all have different colors.

This is an example of what is called a *graph coloring problem*: given a graph G ,

assign colors to each node such that adjacent nodes have different colors. A color

assignment with this property is called a *valid coloring* of the graph —a “*coloring*,”

for short. A graph G is k -colorable if it has a coloring that uses at most k colors.

Definition 9.6.1. The minimum value of k for which a graph, G , has a valid coloring is called its *chromatic number*, $\chi(G)$.

In general, trying to figure out if you can color a graph with a fixed number of colors can take a long time. It's a classic example of a problem for which no fast algorithms are known. In fact, it is easy to check if a coloring works, but it seems

really hard to find it. If you figure out how, then you can get a \$1 million Clay prize.)

5.3.2

9.6.1 Degree-bounded Coloring

There are some simple graph properties that give useful upper bounds on the chromatic number. For example, if we have an upper bound k on the degrees of all the vertices in a graph,

then we can easily find a coloring with only one more color than the degree bound.

Theorem 9.6.2. A graph with maximum degree at most k is $(k + 1)$ -colorable.

The natural way to try to prove this theorem is to use induction on k . Unfortunately, if you try induction on k , it will lead to disaster. It is not that

Unfortunately, this approach leads to disaster.

~~your~~ week

it is impossible, just that it is extremely painful and would ruin you if you tried

~~When you encounter such a disaster when using induction on graphs, it is usually best to change what you are~~

~~inducting on. In graphs, some good choices are n , the number of nodes, or e , the~~

number of edges.

of Theorem 9.6.2

Proof We use induction on the number of vertices in the graph, which we denote

by n . Let $P(n)$ be the proposition that an n -vertex graph with maximum degree at

most k is $(k + 1)$ -colorable.

Base case: ($n = 1$) A 1-vertex graph has maximum degree 0 and is 1-colorable,

so $P(1)$ is true.

Inductive step: Now assume that $P(n)$ is true, and let G be an $(n + 1)$ -vertex

graph with maximum degree at most k . Remove a vertex v (and all edges incident

to it), leaving an n -vertex subgraph, H . The maximum degree of H is at most k ,

and so H is $(k + 1)$ -colorable by our assumption $P(n)$. Now add back vertex v . We

~~pick from the set of $k+1$ colors) that is~~

can assign v a color different from all its adjacent vertices, since there are at most

~~adjacent to v and so at least ~~one~~ one of the~~

~~k adjacent vertices and $k + 1$ colors are available. Therefore, G is $(k + 1)$ -colorable.~~

~~is still~~

~~If we ever will never go in past this last~~

~~students to prove this~~

This completes the inductive step, and the theorem follows by induction. ■

Sometimes $k + 1$ colors is the best you can do. For example, in the complete

graph, K_n , every one of its n vertices is adjacent to all the others, so all n must

be assigned different colors. Of course n colors is also enough, so $\chi(K_n) = n$.

~~In this case, every node has degree $n - 1$ and so this is an example where Theorem 9.6.2 gives the best possible bound. This By a~~

~~means that Theorem 9.6.2 gives the best possible bound for any graph with similar argument, we can show~~

degree bounded by k that has K_{k+1} as a subgraph.

EDITING NOTE:

The complete graph, K_n , is also called a size n clique, just like a clique of friends, where nodes represent the people and an edge represents the friendship relationship.⁶

For example,

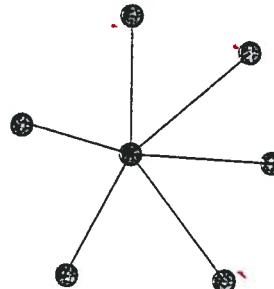
But sometimes $k + 1$ colors is far from the best that you can do. ~~Here's an example~~

⁶ When speaking of friends, clique is usually pronounced similar to click. However, for some reason, graph theorists think that the word clique rhymes with geek.

The n -node star graph shown in Figure 5T has maximum degree $n-1$ but can be colored using just 2 colors.

420

CHAPTER 9. SIMPLE GRAPHS



Example of an n -node star graph for $n = 7$.

A
Figure 5T: A 7-node star graph.

In the n -node star graph, the maximum degree is $n-1$, but the star only needs 2 colors!

5,3,3

9.6.2 Why coloring?

One reason coloring problems come all the time is because scheduling conflicts

are so common. For example, at Akamai, a new version of software is deployed

~65,000 servers

over each of ~~2000~~ servers every few days. The updates cannot be done at the

same time since the servers need to be taken down in order to deploy the software.

Also, the servers cannot be handled one at a time, since it would take forever to

update them all (each one takes about an hour). Moreover, certain pairs of servers

cannot be taken down at the same time since they have common critical functions.

This problem was eventually solved by making a ~~2000~~^{65,000} node conflict graph and

coloring it with 8 colors – so only 8 waves of install are needed! Another example

comes from the need to assign frequencies to radio stations. If two stations have an

overlap in their broadcast area, they can't be given the same frequency. Frequencies are precious and expensive, so you want to minimize the number handed out.

This amounts to finding the minimum coloring for a graph whose vertices are the

stations and whose edges ~~are between~~^{connect} stations with overlapping areas.

Coloring also comes up in allocating registers for program variables. While a

variable is in use, its value needs to be saved in a register. ~~But~~^{so} registers can often be

reused for different variables. ~~But~~^{but} two variables need different registers if they are

referenced during overlapping intervals of program execution. So register alloca-

tion is the coloring problem for a graph whose vertices are the variables; vertices

are adjacent if their intervals overlap, and the colors are registers. *once again, the goal is to minimize the number of colors needed to color the graph.*

Finally, there's the famous map coloring problem stated in Proposition 1.5.4.

The question is how many colors are needed to color a map so that adjacent ter-

Daniel: please
check this

ritories get different colors? This is the same as the number of colors needed to color a graph that can be drawn in the plane without edges crossing. A proof that four colors are enough for ~~the~~ planar graphs was acclaimed when it was discovered about thirty years ago. Implicit in that proof was a 4-coloring procedure that takes time proportional to the number of vertices in the graph (countries in the map).

Surprisingly,
~~On the other hand~~, it's another of those million dollar prize questions to find an efficient procedure to tell if a planar graph really *needs* four colors or if three will actually do the job. *(It's* always easy to tell if an *arbitrary* graph is 2-colorable.)
This chapter,
as we show in Section 9.7. Later in Chapter 11, we'll develop enough planar graph theory to present an easy proof at ~~least~~ *all* that planar graphs are 5-colorable.

9.7. BIPARTITE MATCHINGS

9.6.3 Problems

Class Problems

Homework Problems

Exam Problems

DAVID:

The material in the following section is ~~INSERT H~~ and goes to page 371.

~~9.7~~

9.7 Bipartite Matchings

5.2.2

~~9.7.1~~

9.7.1 Bipartite Graphs

now a subsection

now a subsubsection (no #)

There were two kinds of vertices in the "Sex in America" graph —males and females,

males, and edges only went between the two kinds. Graphs like this come up so

frequently they have earned a special name —they are called *bipartite graphs*.

Definition 9.7.1. A *bipartite graph* is a graph together with a partition of its vertices

into two sets, L and R , such that every edge is incident to a vertex in L and to a

vertex in R .

So every bipartite graph looks something like ~~this~~ ^{9.1.} the graph in Figure 5-~~5~~

— ~~INSERT J~~ goes here —

The Problem \Leftarrow subsubsection (no #)

The bipartite matching problem is related to the ~~satistics problem~~ sex in America problem that we just studied, only now the goal is to get everyone happily married, or at least to pair up every man with a women that he likes, and vice-versa. ~~As you can imagine, this~~ As you might imagine, this is not possible for a variety of reasons, not the least of which is the fact that there are more women ~~than~~ in America than men. So, it is simply not possible to ~~match every woman with a unique man~~ pair up every woman with a man ~~elsewhere so~~ so that every man is paired with at most one woman.

But what about getting a mate for every man so that every woman is paired with at most one man? Is it possible to do this so that each man is paired with a woman that he likes? The answer, ^{of course,} depends on the bipartite graph that represents who likes who, but the good news is that ~~it is~~ it is possible to characterize ~~the answer to this question in terms of natural~~

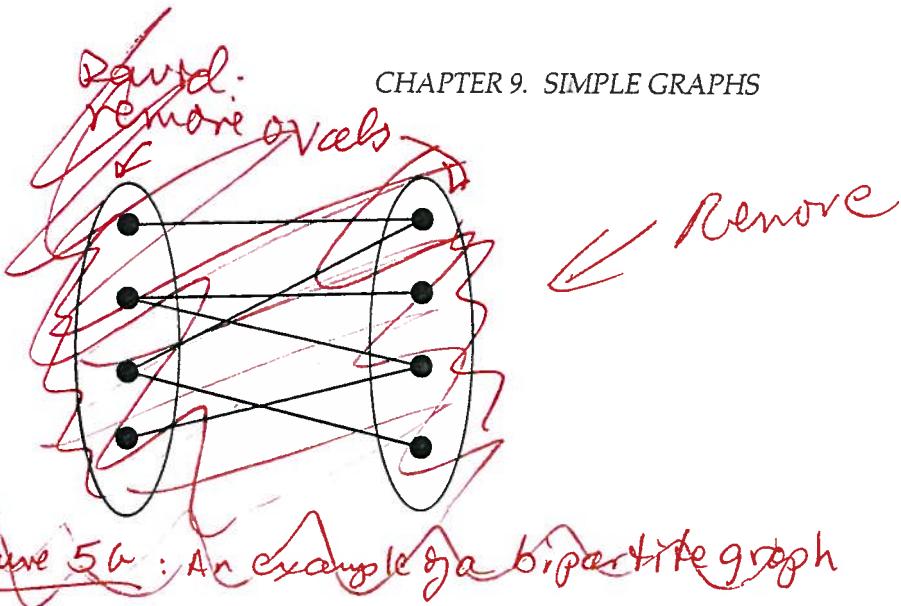
~~the who-likes-who graph~~

find natural properties of the who-likes-who graph that, determine the answer to this question completely.

In general, suppose we have a ~~collection~~ ^{that} set of men and an equal-sized or larger set of women, and there is a graph with an edge between a man and a woman if the man likes the woman. (Note that in this scenario, the "likes" relationship need not be symmetric since for the time being, we will only worry about finding a mate for a man that he likes. Later, we will consider the "likes" relationship from the female perspective as well.) For example, we might obtain the graph in Figure 5J.

read too much into this example.
~~all this too~~
~~take offense~~
~~seriously.~~)

I By the way, we ~~are not taking any positions~~ do not mean to imply ~~any~~ that ~~positions~~ mates or pairs should do should not be of a heterosexual nature. ~~nor~~ Nor do we mean to imply that men should get their choice instead of women. It's just that with bipartite graphs, the edges only connect ~~men to women~~ male nodes to female nodes and there are fewer men in America. ~~with that said~~ so please don't ~~sue us~~.



Now we can immediately see how to color a bipartite graph using only two colors: let all the L vertices be black and all the R vertices be white. Conversely, if a graph is 2-colorable, then it is bipartite with L being the vertices of one color and R the vertices of the other color. In other words,

"bipartite" is a synonym for "2-colorable."

The following Lemma gives another useful characterization of bipartite graphs.

Theorem 9.7.2. *A graph is bipartite iff it has no odd-length cycle.*

The proof of Theorem 9.7.2 is left to Problem ??.

~~David! Throughout this section, replace "girl" with "woman", "boy" with "man", "girls" with "women", "boys" with "men", etc. Remove "title"~~

9.7. BIPARTITE MATCHINGS

425

9.7.2 Bipartite Matchings

~~INSERT~~

The *bipartite matching* problem resembles the stable Marriage Problem in that it concerns a set of girls and a set of at least as many boys. There are no preference lists, but each girl does have some boys she likes and others she does not like. In the bipartite matching problem, we ask whether every girl can be paired up with a boy that she likes. Any particular matching problem can be specified by a bipartite graph with a vertex for each girl, a vertex for each boy, and an edge between a boy and a girl iff the girl likes the boy. For example, we might obtain the following graph:

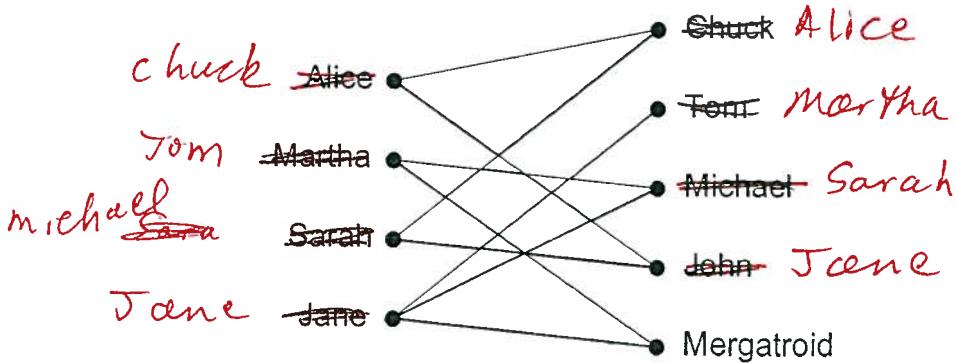


Figure 5J: A graph where an edge between a man and a woman denotes that the man likes the woman.

In this ~~problem~~, ~~a matching will mean a way of assigning every girl to a boy so that different girls are assigned to different boys, and a girl is always assigned to a boy she likes~~

likes. For example, ~~here~~ one possible matching for the ~~girls~~ men is shown in Figure ~~5K~~

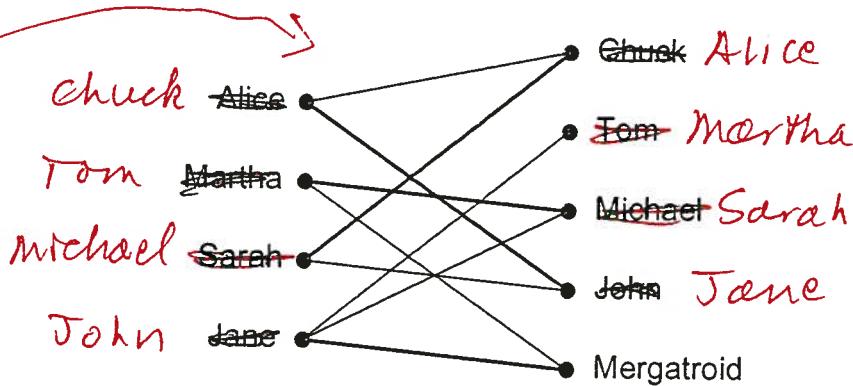


Figure 5K: One possible matching for the men.

Hall's Matching Theorem ~~states~~ gives necessary and sufficient conditions for the ex-

A famous result known as

istence of a matching in a bipartite graph. It turns out to be a remarkably useful

mathematical tool.

Subsubsection
(no #)

9.7.3 The Matching Condition

man-likes-women

We'll state and prove Hall's Theorem using ~~girl-likes-boy~~ terminology. Define the women

set of ~~boys~~ liked by a given set of ~~girls~~ to consist of all ~~boys~~ liked by at least one of

men women Tom John martha
~~men~~ ~~women~~ ~~Tom~~ ~~John~~ ~~martha~~
those ~~girls~~. For example, the set of ~~boys~~ liked by ~~Martha~~ and ~~Jane~~ consists of ~~Tom~~,
Sarah, ~~Michael~~, and Mergatroid. For us to have any chance at all of matching up the ~~girls~~,

the following *matching condition* must hold:

9.7. BIPARTITE MATCHINGS

427

men *women*
Every subset of ~~girls~~ likes at least as large a set of ~~boys~~.

set of 4 men *women*

For example, we can not find a matching if some ~~4 girls~~ like only 3 ~~boys~~. Hall's

Theorem says that this necessary condition is actually sufficient; if the matching condition holds, then a matching exists.

men M women W

Theorem 9.7.3. A matching for a set of ~~girls~~ with a set of ~~boys~~ can be found if and only if the matching condition holds.

Proof. First, let's suppose that a matching exists and show that the matching con-

dition holds. Consider an arbitrary subset of ~~girls~~. Each ~~girl~~ likes at least the ~~boys~~ *she* ~~is~~ is matched with. Therefore, every subset of ~~girls~~ likes at least as large a set of ~~boys~~. Thus, the matching condition holds.

Next, let's suppose that the matching condition holds and show that a matching

exists. We use strong induction on ~~M~~, the number of ~~girls~~.

m m

Base Case: ($|M| = 1$) If $|M| = 1$, then the matching condition implies that the

man women
alone ~~girl~~ likes at least one ~~boy~~, and so a matching exists.

m

Inductive Step: Now suppose that $|M| \geq 2$. There are two cases:

footnote

1) men
~~girls~~

women

Case 1: Every proper subset of girls likes a *strictly larger* set of boys. In this case, we

man women he

have some latitude: we pair an arbitrary girl with a boy he likes and send

them both away. The matching condition still holds for the remaining boys

women since we have removed ~~at most~~ ^{only} one woman, and ~~girls~~, so we can match the rest of the girls by induction.

men

men

women ~~girls~~ $\forall w \in W$.

Case 2: Some proper subset of ~~girls~~ $X \subset G$ likes an *equal-size* set of boys ~~boys~~ $Y \subseteq B$.

men in M women in W

We match the girls in X with the boys in Y by induction and send them all

men

away. We can also match the rest of the girls by induction if we show that

men women

the matching condition holds for the remaining boys and girls. To check the

matching condition for the remaining people, consider an arbitrary subset of

men m

the remaining ~~girls~~ $X' \subseteq (G - X)$, and let Y' be the set of remaining boys

Women

that they like. We must show that $|X'| \leq |Y'|$. Originally, the combined set

men women

of ~~girls~~ $X \cup X'$ liked the set of boys $Y \cup Y'$. So, by the matching condition,

we know:

$$|X \cup X'| \leq |Y \cup Y'|$$

men

We sent away $|X|$ girls from the set on the left (leaving X') and sent away

I recall that a subset A of B is proper if $A \neq B$.

women

an equal number of ~~boys~~^{women} from the set on the right (leaving Y'). Therefore, it

must be that $|X'| \leq |Y'|$ as claimed.

in both cases, there is

men,

So there is in any case a matching for the ~~girls~~^{women}, which completes the proof of

the Inductive step. The theorem follows by induction. ■

9.7.3

The proof of ~~Hall's~~^{Hall} theorem gives an algorithm for finding a matching in a bipar-

tite graph, albeit not a very efficient one. However, efficient algorithms for finding

a matching in a bipartite graph do exist. Thus, if a problem can be reduced to

finding a matching, the problem is essentially solved from a computational per-

spective.

↙ subsubsection (no #)

9.7.4 A Formal Statement

9.7.3

Let's restate ~~Hall's~~^{Hall} Theorem in abstract terms so that you'll not always be con-

men

women

demned to saying, "Now this group of ~~little girls~~^{men} likes at least as many ~~little boys~~^{women}..."

Definition 5&K!

↑ A matching in a graph, G , is a set of edges such that no two edges in the set

share a vertex. A matching is said to cover a set, L , of vertices iff each vertex in L

A matching is said to be perfect; if every node in the graph is incident to an edge in the matching.

M-IV

430

CHAPTER 9. SIMPLE GRAPHS

has an edge of the matching incident to it. In any graph, the set $N(S)$, of neighbors of some set, S , of vertices is the set of all vertices adjacent to some vertex in S . That

is,

$$N(S) := \{r \mid s—r \text{ is an edge for some } s \in S\}.$$

S is called a *bottleneck* if

$$|S| > |N(S)|.$$

Theorem 9.7.4 (Hall's Theorem). Let G be a bipartite graph with vertex partition L, R .

There is matching in G that covers L iff no subset of L is a bottleneck.

~~sufficient~~
 An Easy Matching Condition ← subsubsection (no #)

men

The bipartite matching condition requires that every subset of ~~girls~~ has a certain

property. In general, verifying that every subset has some property, even if it's easy

to check any particular subset for the property, quickly becomes overwhelming

because the number of subsets of even relatively small sets is enormous —over a

⁷An equivalent definition of $N(S)$ uses relational notation. $N(S)$ is simply the image SR of S under the adjacency relation, R , on vertices of the graph.

billion subsets for a set of size 30. However, there is a simple property of vertex

degrees in a bipartite graph that guarantees a match and is very easy to check.

Namely call a bipartite graph *degree-constrained* if vertex degrees on the left are at

least as large as those on the right. More precisely,

Definition 9.7.5. A bipartite graph G with vertex partition L, R is *degree-constrained*

if $\deg(l) \geq \deg(r)$ for every $l \in L$ and $r \in R$.

~~It turns out that~~

~~Now we can always find a matching in a degree-constrained bipartite graph.~~

~~The following lemma implies that~~

Lemma 9.7.6. Every degree constrained bipartite graph satisfies the matching condition.

~~Let consider a bipartite~~

Proof. Let S be any set of vertices in L . The number of edges incident to vertices

in S is exactly the sum of the degrees of the vertices in S . Each of these edges is

incident to a vertex in $N(S)$ by definition of $N(S)$. So the sum of the degrees of

the vertices in $N(S)$ is at least as large as the sum for S . But since the degree of

every vertex in $N(S)$ is at most as large as the degree of every vertex in S , there

would have to be at least as many terms in the sum for $N(S)$ as in the sum for S .

So there have to be at least as many vertices in $N(S)$ as in S , proving that S is not a

~~bottleneck~~

~~— INSERT N goes here —~~

Corollary 5.1

For example, the graph in Figure 5.15 is degree constrained since every node on the left is adjacent to at least two nodes on the right while every node on the right is incident to at most two nodes on the left. ~~regular~~

Regular graphs often arise in practice provide a large class of graphs that, ^{often} arise in practice that are degree constrained.

Definition 5P: A graph is said to be regular, if every node has the same degree.

The following lemma implies that every degree constrained bipartite graph has a matching. ~~that covers~~

Lemma 9.76: Let G be a bipartite graph with vertex partition L, R ~~where $|L| \leq |R|$~~ where $|L| \leq |R|$. If G is degree-constrained, then no subset of L is a bottleneck.

Corollary 5L: Let G be a bipartite graph with vertex partition L, R where $|L| \leq |R|$. If G is degree-constrained, then there is a matching that covers L .

Proof: Combine Lemma 9.7.6 with

Theorem 9.7.4. \blacksquare
we can now prove:

Theorem

Corollary 5M: Every regular bipartite

graph has a perfect matching.

Proof: Let G be a bipartite graph with vertex partition L, R where $|L| \leq |R|$. Since every node has the same degree, and since $|L| \leq |R|$, regular graphs are degree-constrained, we know by Corollary 5L that there must be a matching in G that covers L . Since $|L| \leq |R|$, this matching also covers R . Since G is regular, we also know that $|L| = |R|$ and thus the matching must also cover R . This means that every node in G is incident to an edge in the matching and thus G has a perfect matching. \blacksquare

Corollary 5M with poore arises in numerous applications in later chapters where we need to assert that certain graphs have perfect matchings

~~Corollary 5A.5a's~~

Theorem 5M is ~~so~~ a surprisingly useful tool in graph theory. For example, we will use it in chapter 6 when we are figuring out how to route data paths in a Series Network.

Everything is crossed out on
this page.

CHAPTER 9. SIMPLE GRAPHS

432

bottleneck. So there are no bottlenecks, proving that the degree-constrained graph

satisfies the matching condition.

Of course being degree-constrained is a very strong property, and lots of graphs

that aren't degree-constrained have matchings. But we'll see examples of degree-

constrained graphs come up naturally in some later applications.

9.7.5 Problems

Class Problems

Exam Problems

Homework Problems

Chapter 9

Definition 5X:

A matching is said to be perfect if every node is incident to an edge in the

Corollary 5L: Let $G = (V, E)$ be a degree constrained bipartite graph, where V is partitioned into L and R . Then there is a perfect matching for G that covers L .

Corollary 5M: There is a perfect matching that covers every node in any regular bipartite graph.

Proof: If every node has the same degree, then $|L| = |R|$ and a matching that covers L also covers R . Such a matching must exist by Lemma 9.7.6. \blacksquare

THIS IS THE END OF INSERT H.